

```
In [1]: 1 import numpy as np  
2 import pandas as pd  
3 import matplotlib.pyplot as plt
```

```
In [2]: 1 df = pd.read_csv('laptop_data.csv')
```

```
In [3]: 1 df.head()
```

Out[3]:

| | Unnamed: 0 | Company | TypeName | Inches | ScreenResolution | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price |
|---|---------------|---------|-----------|--------|---------------------------------------|----------------------------------|------|---------------------------|---------------------------------------|-------|--------|-------------|
| 0 | 0 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 2.3GHz | 8GB | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37kg | 71378.6832 |
| 1 | 1 | Apple | Ultrabook | 13.3 | 1440x900 | Intel Core i5 1.8GHz | 8GB | 128GB Flash Storage | Intel HD Graphics 6000 | macOS | 1.34kg | 47895.5232 |
| 2 | 2 | HP | Notebook | 15.6 | Full HD 1920x1080 | Intel Core i5 7200U 2.5GHz | 8GB | 256GB SSD | Intel HD Graphics 620 | No OS | 1.86kg | 30636.0000 |
| 3 | 3 | Apple | Ultrabook | 15.4 | IPS Panel Retina Display 2880x1800 | Intel Core i7 2.7GHz | 16GB | 512GB SSD | AMD Radeon Pro 455 | macOS | 1.83kg | 135195.3360 |
| 4 | 4 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 3.1GHz | 8GB | 256GB SSD | Intel Iris Plus Graphics 650 | macOS | 1.37kg | 96095.8080 |

In [4]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1303 entries, 0 to 1302
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        1303 non-null    int64  
 1   Company          1303 non-null    object  
 2   TypeName         1303 non-null    object  
 3   Inches           1303 non-null    float64 
 4   ScreenResolution 1303 non-null    object  
 5   Cpu              1303 non-null    object  
 6   Ram              1303 non-null    object  
 7   Memory           1303 non-null    object  
 8   Gpu              1303 non-null    object  
 9   OpSys            1303 non-null    object  
 10  Weight           1303 non-null    object  
 11  Price            1303 non-null    float64 
dtypes: float64(2), int64(1), object(9)
memory usage: 122.3+ KB
```

In [5]: 1 df.isnull().sum()

```
Out[5]: Unnamed: 0      0
Company          0
TypeName         0
Inches           0
ScreenResolution 0
Cpu              0
Ram              0
Memory           0
Gpu              0
OpSys            0
Weight           0
Price            0
dtype: int64
```

```
In [6]: 1 df.drop(columns=['Unnamed: 0'], inplace=True)
```

```
In [7]: 1 df.head()
```

Out[7]:

| | Company | TypeName | Inches | ScreenResolution | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price |
|---|---------|-----------|--------|------------------------------------|----------------------------|------|---------------------|------------------------------|-------|--------|-------------|
| 0 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 2.3GHz | 8GB | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37kg | 71378.6832 |
| 1 | Apple | Ultrabook | 13.3 | 1440x900 | Intel Core i5 1.8GHz | 8GB | 128GB Flash Storage | Intel HD Graphics 6000 | macOS | 1.34kg | 47895.5232 |
| 2 | HP | Notebook | 15.6 | Full HD 1920x1080 | Intel Core i5 7200U 2.5GHz | 8GB | 256GB SSD | Intel HD Graphics 620 | No OS | 1.86kg | 30636.0000 |
| 3 | Apple | Ultrabook | 15.4 | IPS Panel Retina Display 2880x1800 | Intel Core i7 2.7GHz | 16GB | 512GB SSD | AMD Radeon Pro 455 | macOS | 1.83kg | 135195.3360 |
| 4 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 3.1GHz | 8GB | 256GB SSD | Intel Iris Plus Graphics 650 | macOS | 1.37kg | 96095.8080 |

```
In [8]: 1 df['Ram'] = df['Ram'].str.replace('GB', '')
2 df['Weight'] = df['Weight'].str.replace('kg', '')
```

```
In [9]: 1 df.head()
```

Out[9]:

| | Company | TypeName | Inches | ScreenResolution | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price |
|---|---------|-----------|--------|------------------------------------|----------------------------|-----|---------------------|------------------------------|-------|--------|-------------|
| 0 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 2.3GHz | 8 | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 |
| 1 | Apple | Ultrabook | 13.3 | 1440x900 | Intel Core i5 1.8GHz | 8 | 128GB Flash Storage | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 |
| 2 | HP | Notebook | 15.6 | Full HD 1920x1080 | Intel Core i5 7200U 2.5GHz | 8 | 256GB SSD | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 |
| 3 | Apple | Ultrabook | 15.4 | IPS Panel Retina Display 2880x1800 | Intel Core i7 2.7GHz | 16 | 512GB SSD | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 |
| 4 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 3.1GHz | 8 | 256GB SSD | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 |

```
In [10]: 1 df['Ram'] = df['Ram'].astype('int32')
2 df['Weight'] = df['Weight'].astype('float32')
```

```
In [11]: 1 df.info()
```

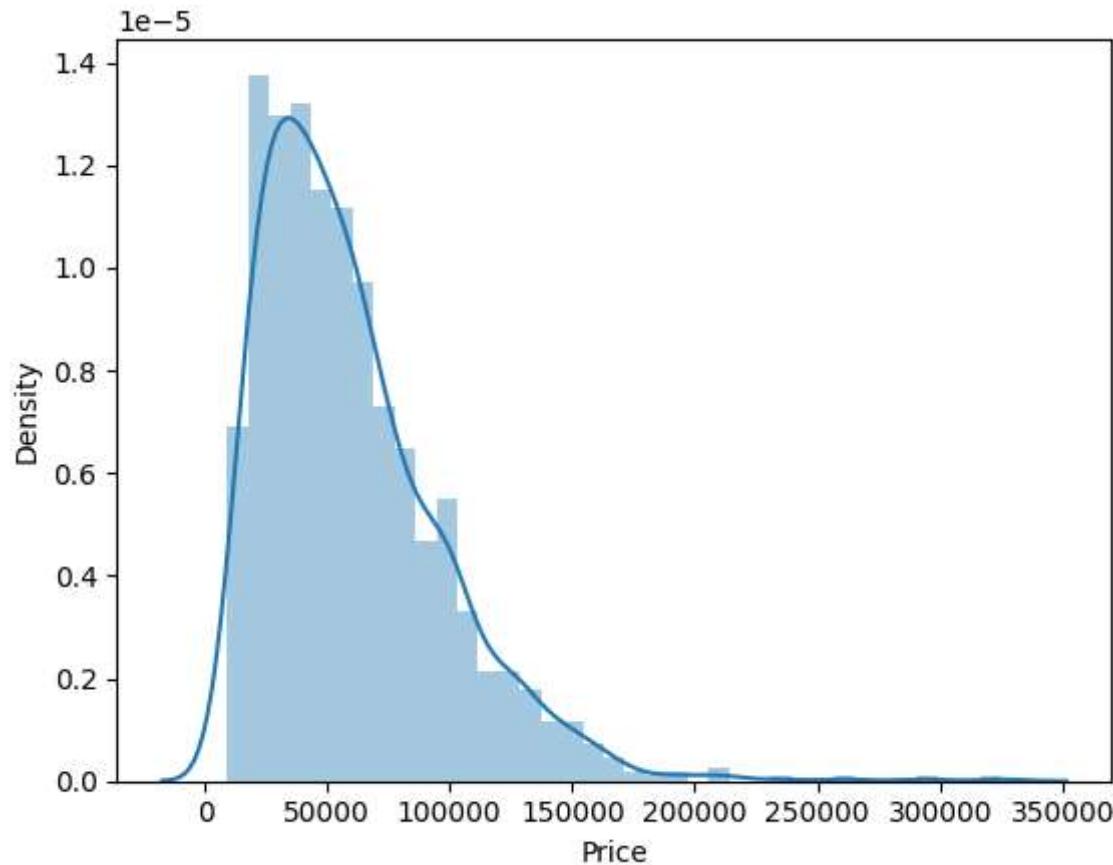
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1303 entries, 0 to 1302
Data columns (total 11 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Company          1303 non-null    object  
 1   TypeName         1303 non-null    object  
 2   Inches           1303 non-null    float64 
 3   ScreenResolution 1303 non-null    object  
 4   Cpu              1303 non-null    object  
 5   Ram              1303 non-null    int32   
 6   Memory           1303 non-null    object  
 7   Gpu              1303 non-null    object  
 8   OpSys            1303 non-null    object  
 9   Weight            1303 non-null    float32 
 10  Price             1303 non-null    float64 
dtypes: float32(1), float64(2), int32(1), object(7)
memory usage: 101.9+ KB
```

```
In [12]: 1 import seaborn as sns
```

```
In [13]: 1 sns.distplot(df["Price"])
```

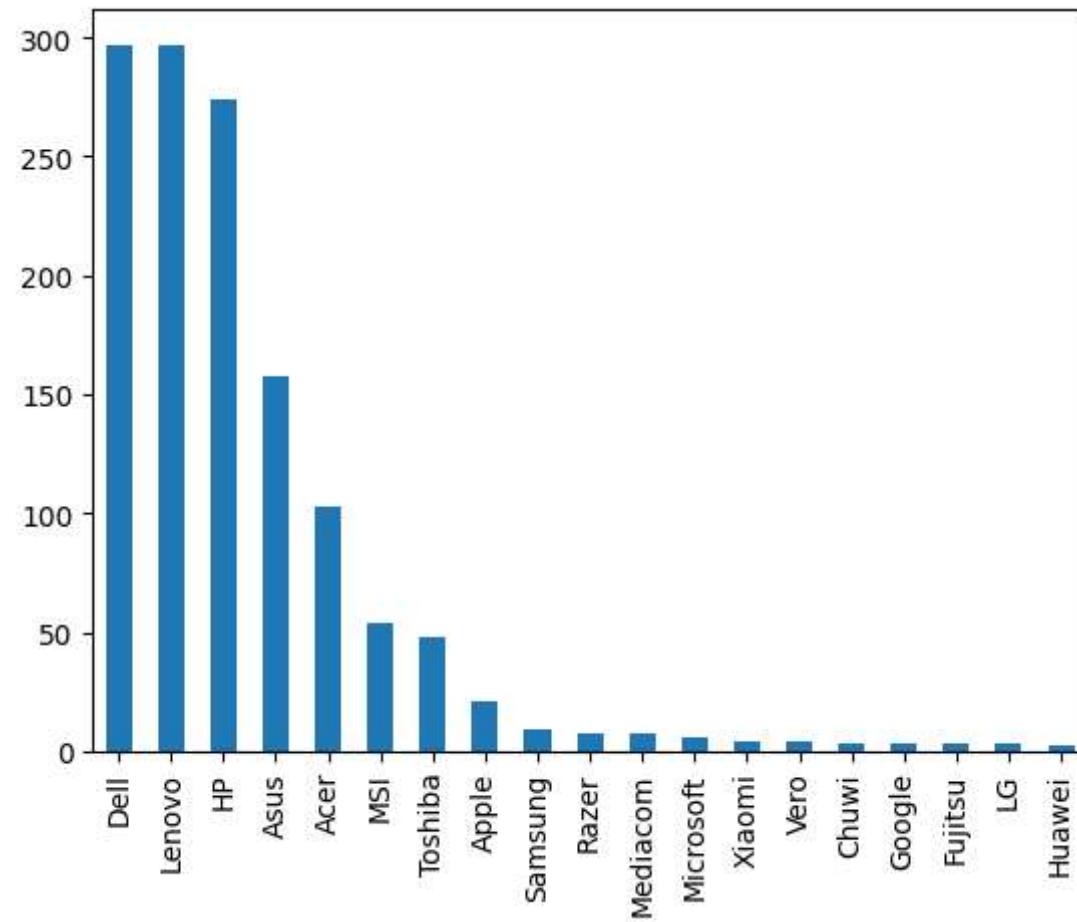
C:\Users\sanket parekh\AppData\Local\Temp\ipykernel_35592\941010651.py:1: UserWarning:
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
Please adapt your code to use either `displot` (a figure-level function with
similar flexibility) or `histplot` (an axes-level function for histograms).
For a guide to updating your code to use the new functions, please see
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751> (<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>)
sns.distplot(df["Price"])

```
Out[13]: <AxesSubplot: xlabel='Price', ylabel='Density'>
```



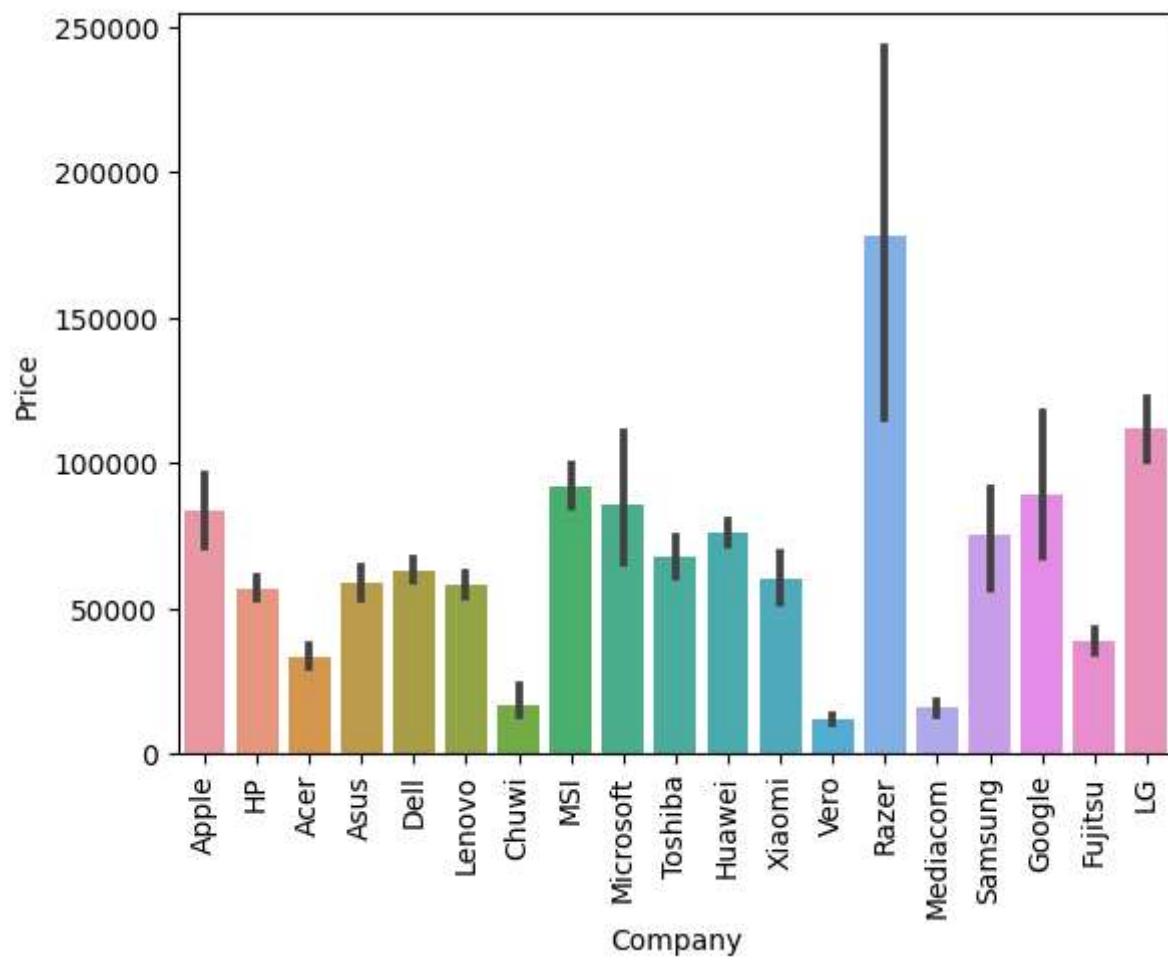
```
In [14]: 1 df['Company'].value_counts().plot(kind='bar')
```

```
Out[14]: <AxesSubplot: >
```



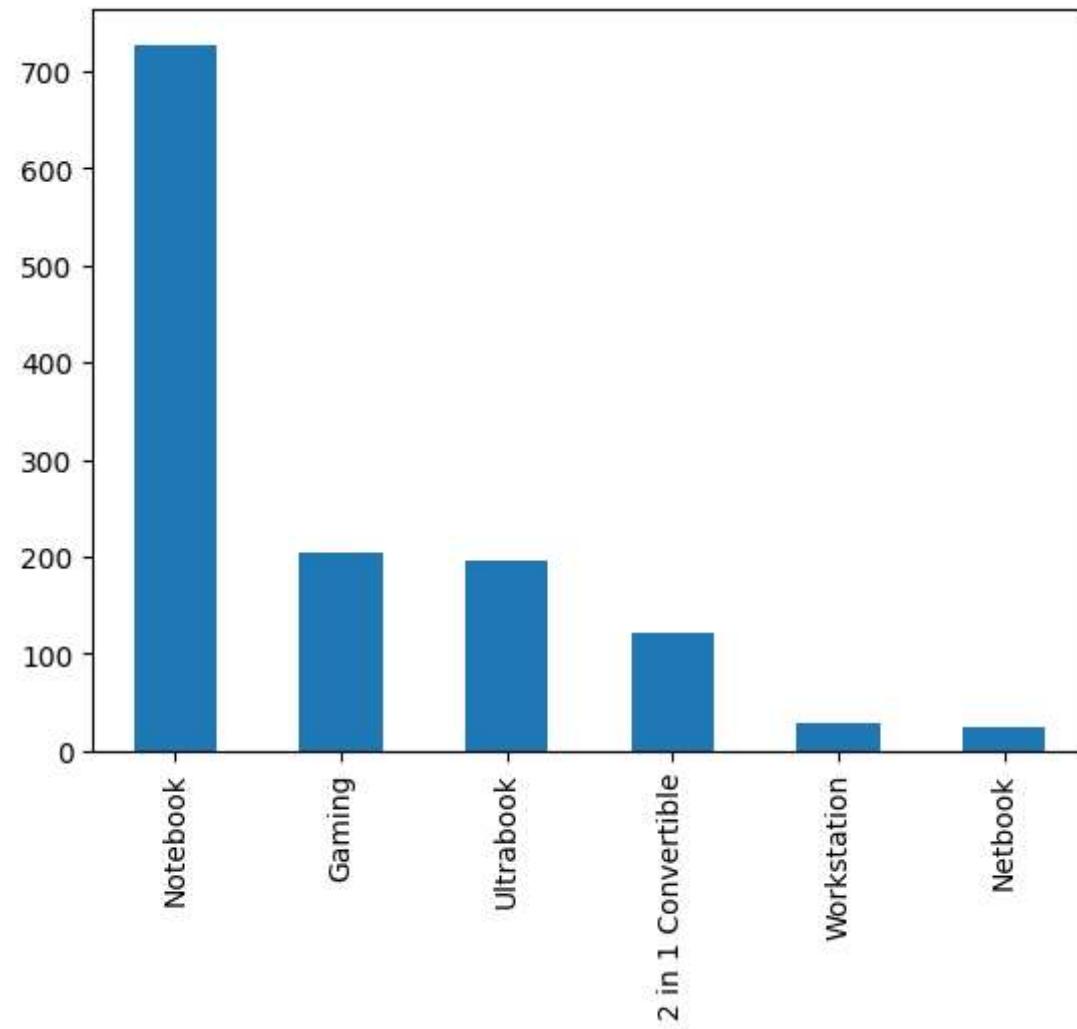
In [23]:

```
1 sns.barplot(x=df['Company'],y=df['Price'])
2 plt.xticks(rotation='vertical')
3 plt.show()
```



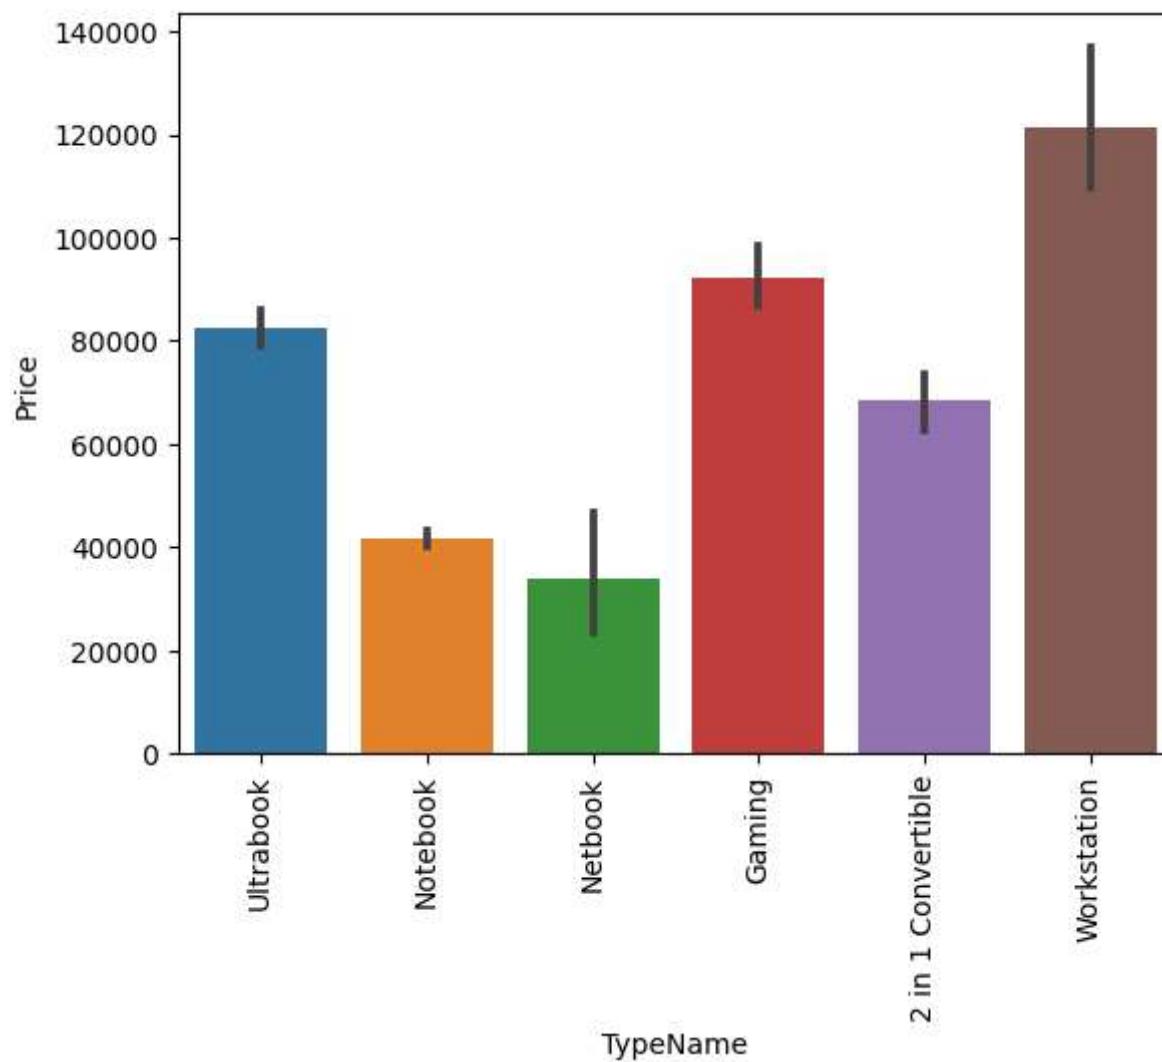
```
In [24]: 1 df['TypeName'].value_counts().plot(kind='bar')
```

```
Out[24]: <AxesSubplot: >
```



In [25]:

```
1 sns.barplot(x=df['TypeName'],y=df['Price'])
2 plt.xticks(rotation='vertical')
3 plt.show()
```



```
In [26]: 1 sns.distplot(df['Inches'])
```

C:\Users\sanket parekh\AppData\Local\Temp\ipykernel_35592\1439577752.py:1: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

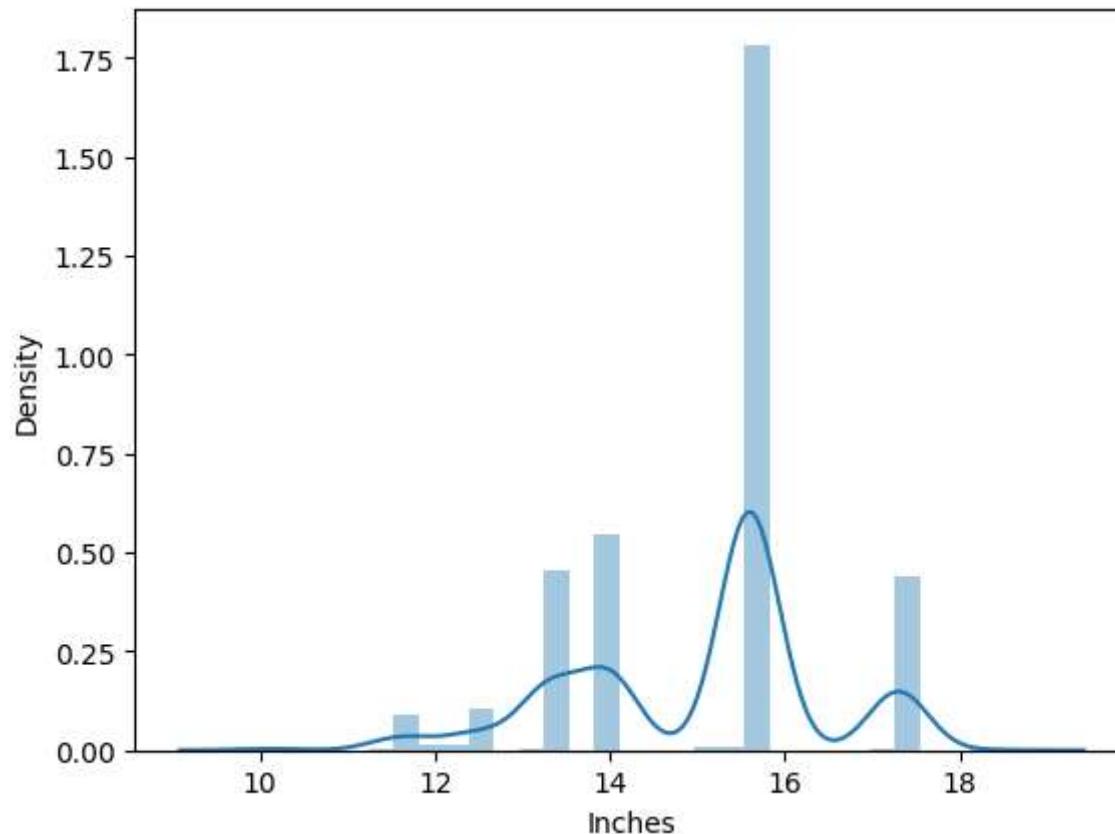
Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751> (<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>)

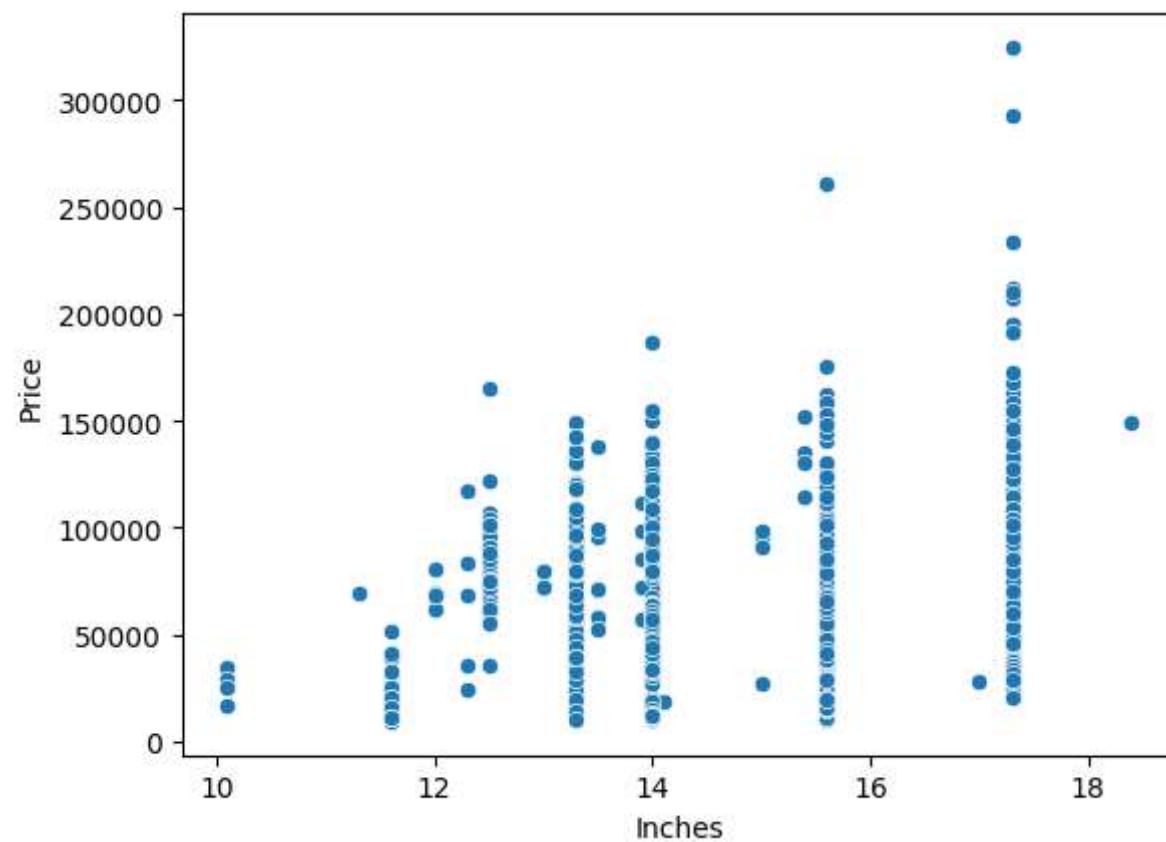
```
sns.distplot(df['Inches'])
```

Out[26]: <AxesSubplot: xlabel='Inches', ylabel='Density'>



```
In [27]: 1 sns.scatterplot(x=df['Inches'],y=df['Price'])
```

```
Out[27]: <AxesSubplot: xlabel='Inches', ylabel='Price'>
```



In [28]: 1 df['ScreenResolution'].value_counts()

Out[28]:

| | |
|---|-----|
| Full HD 1920x1080 | 507 |
| 1366x768 | 281 |
| IPS Panel Full HD 1920x1080 | 230 |
| IPS Panel Full HD / Touchscreen 1920x1080 | 53 |
| Full HD / Touchscreen 1920x1080 | 47 |
| 1600x900 | 23 |
| Touchscreen 1366x768 | 16 |
| Quad HD+ / Touchscreen 3200x1800 | 15 |
| IPS Panel 4K Ultra HD 3840x2160 | 12 |
| IPS Panel 4K Ultra HD / Touchscreen 3840x2160 | 11 |
| 4K Ultra HD / Touchscreen 3840x2160 | 10 |
| 4K Ultra HD 3840x2160 | 7 |
| Touchscreen 2560x1440 | 7 |
| IPS Panel 1366x768 | 7 |
| IPS Panel Quad HD+ / Touchscreen 3200x1800 | 6 |
| IPS Panel Retina Display 2560x1600 | 6 |
| IPS Panel Retina Display 2304x1440 | 6 |
| Touchscreen 2256x1504 | 6 |
| IPS Panel Touchscreen 2560x1440 | 5 |
| IPS Panel Retina Display 2880x1800 | 4 |
| IPS Panel Touchscreen 1920x1200 | 4 |
| 1440x900 | 4 |
| IPS Panel 2560x1440 | 4 |
| IPS Panel Quad HD+ 2560x1440 | 3 |
| Quad HD+ 3200x1800 | 3 |
| 1920x1080 | 3 |
| Touchscreen 2400x1600 | 3 |
| 2560x1440 | 3 |
| IPS Panel Touchscreen 1366x768 | 3 |
| IPS Panel Touchscreen / 4K Ultra HD 3840x2160 | 2 |
| IPS Panel Full HD 2160x1440 | 2 |
| IPS Panel Quad HD+ 3200x1800 | 2 |
| IPS Panel Retina Display 2736x1824 | 1 |
| IPS Panel Full HD 1920x1200 | 1 |
| IPS Panel Full HD 2560x1440 | 1 |
| IPS Panel Full HD 1366x768 | 1 |
| Touchscreen / Full HD 1920x1080 | 1 |
| Touchscreen / Quad HD+ 3200x1800 | 1 |
| Touchscreen / 4K Ultra HD 3840x2160 | 1 |
| IPS Panel Touchscreen 2400x1600 | 1 |

Name: ScreenResolution, dtype: int64

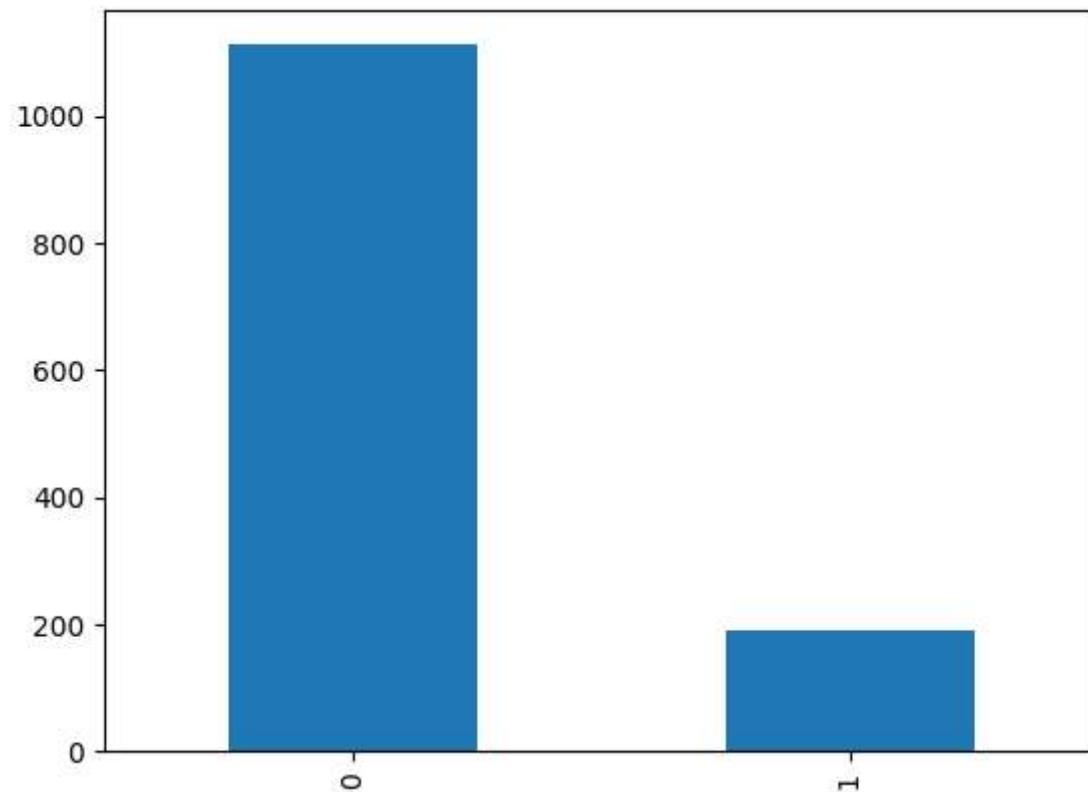
```
In [32]: 1 df['Touchscreen'] = df['ScreenResolution'].apply(lambda x:1 if 'Touchscreen' in x else 0)
2 df.sample(5)
```

Out[32]:

| | Company | TypeName | Inches | ScreenResolution | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscreen |
|------|----------|-----------|--------|---------------------------------------|---|-----|------------------------------|-------------------------------------|---------------|--------|------------|-------------|
| 533 | Mediacom | Notebook | 13.3 | IPS Panel Full HD 1920x1080 | Intel Celeron Quad Core N3450 1.1GHz | 4 | 32GB SSD | Intel HD Graphics 500 | Windows 10 | 1.20 | 19660.3200 | 0 |
| 1187 | Acer | Notebook | 15.6 | Full HD 1920x1080 | Intel Core i5 7200U 2.5GHz | 4 | 256GB SSD | Nvidia GeForce 940MX | Windows 10 | 2.23 | 36816.4800 | 0 |
| 522 | Dell | Notebook | 14.0 | Full HD 1920x1080 | Intel Core i5 7300U 2.6GHz | 8 | 256GB SSD | Intel HD Graphics 620 | Windows 10 | 1.64 | 68184.0144 | 0 |
| 768 | Samsung | Ultrabook | 13.3 | Full HD / Touchscreen 1920x1080 | Intel Core i7 7500U 2.7GHz | 8 | 256GB SSD | Intel HD Graphics 620 | Windows 10 | 1.31 | 85194.7200 | 1 |
| 330 | MSI | Gaming | 17.3 | Full HD 1920x1080 | Intel Core i7 7700HQ 2.8GHz | 8 | 256GB SSD + 1TB HDD | Nvidia GeForce GTX 1050 Ti | Windows 10 | 2.70 | 63882.7200 | 0 |

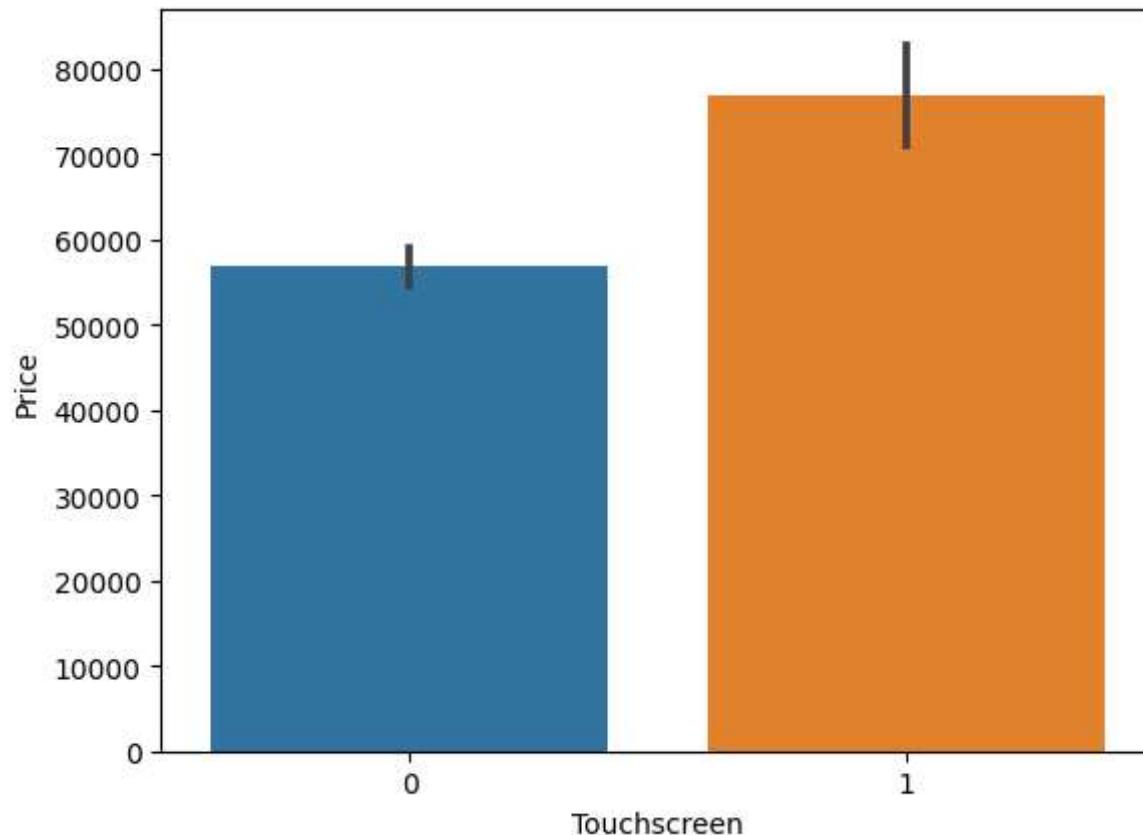
```
In [33]: 1 df['Touchscreen'].value_counts().plot(kind='bar')
```

```
Out[33]: <AxesSubplot: >
```



```
In [34]: 1 sns.barplot(x=df['Touchscreen'],y=df['Price'])
```

```
Out[34]: <AxesSubplot: xlabel='Touchscreen', ylabel='Price'>
```



```
In [35]: 1 df['Ips'] = df['ScreenResolution'].apply(lambda x:1 if 'IPS' in x else 0)
```

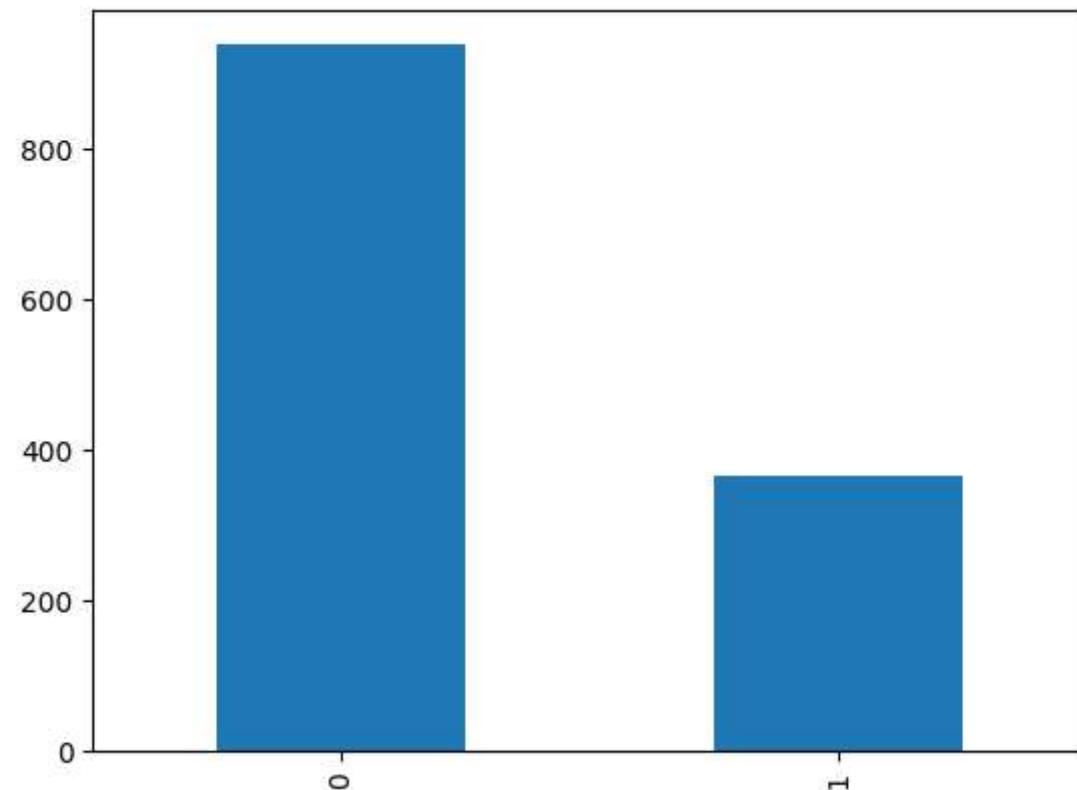
In [36]: 1 df.head()

Out[36]:

| | Company | TypeName | Inches | ScreenResolution | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscreen | Ips |
|---|---------|-----------|--------|------------------------------------|----------------------------|-----|---------------------|------------------------------|-------|--------|-------------|-------------|-----|
| 0 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 2.3GHz | 8 | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | 0 | 1 |
| 1 | Apple | Ultrabook | 13.3 | 1440x900 | Intel Core i5 1.8GHz | 8 | 128GB Flash Storage | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 | 0 | 0 |
| 2 | HP | Notebook | 15.6 | Full HD 1920x1080 | Intel Core i5 7200U 2.5GHz | 8 | 256GB SSD | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 | 0 | 0 |
| 3 | Apple | Ultrabook | 15.4 | IPS Panel Retina Display 2880x1800 | Intel Core i7 2.7GHz | 16 | 512GB SSD | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 | 0 | 1 |
| 4 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 3.1GHz | 8 | 256GB SSD | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 | 0 | 1 |

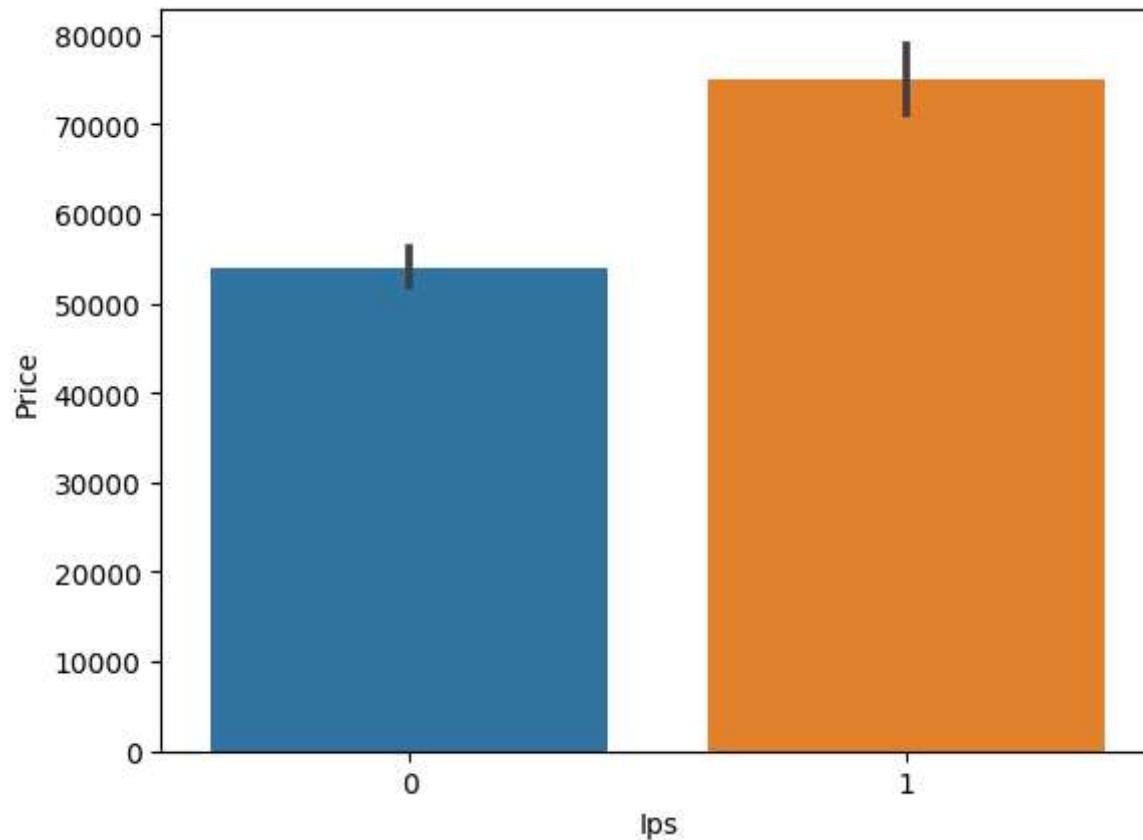
```
In [37]: 1 df['Ips'].value_counts().plot(kind='bar')
```

```
Out[37]: <AxesSubplot: >
```



```
In [38]: 1 sns.barplot(x=df['Ips'],y=df['Price'])
```

```
Out[38]: <AxesSubplot: xlabel='Ips', ylabel='Price'>
```



```
In [39]: 1 new = df['ScreenResolution'].str.split('x',n=1,expand=True)
```

In [40]:

```
1 df['X_res'] = new[0]
2 df['Y_res'] = new[1]
3 df.sample(5)
4
```

Out[40]:

| | | Company | TypeName | Inches | ScreenResolution | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscreen | Ips |
|------|---------|----------|----------|-------------------|----------------------------|-----|-----------|-----------------------|------------|-------|----------|-------|-------------|-----|
| 46 | Lenovo | Notebook | 15.6 | 1366x768 | Intel Core i3 6006U 2GHz | 4 | 128GB SSD | Intel HD Graphics 520 | No OS | 2.20 | 19660.32 | | 0 | 0 |
| 528 | Dell | Notebook | 15.6 | Full HD 1920x1080 | Intel Core i5 7200U 2.5GHz | 4 | 500GB HDD | AMD Radeon R5 M430 | Windows 10 | 2.30 | 31168.80 | | 0 | 0 |
| 862 | Toshiba | Notebook | 13.3 | Full HD 1920x1080 | Intel Core i5 7200U 2.5GHz | 4 | 128GB SSD | Intel HD Graphics 620 | Windows 10 | 1.05 | 68464.80 | | 0 | 0 |
| 1004 | Toshiba | Notebook | 13.3 | Full HD 1920x1080 | Intel Core i5 6200U 2.3GHz | 4 | 128GB SSD | Intel HD Graphics 520 | Windows 10 | 1.20 | 63669.60 | | 0 | 0 |
| 859 | Lenovo | Notebook | 15.6 | Full HD 1920x1080 | Intel Core i3 6006U 2.0GHz | 4 | 1TB HDD | Intel HD Graphics 520 | No OS | 2.20 | 24988.32 | | 0 | 0 |



```
In [41]: 1 df['X_res'] = df['X_res'].str.replace(',','').str.findall(r'(\d+\.\?\d+)').apply(lambda x:x[0])
2 df.head()
3
```

Out[41]:

| | Company | TypeName | Inches | ScreenResolution | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscreen | Ips | X |
|---|---------|-----------|--------|--|-------------------------------------|-----|---------------------------|---------------------------------------|-------|--------|-------------|-------------|-----|---|
| 0 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 2.3GHz | 8 | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | 0 | 1 | |
| 1 | Apple | Ultrabook | 13.3 | 1440x900 | Intel Core i5 1.8GHz | 8 | 128GB Flash Storage | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 | 0 | 0 | |
| 2 | HP | Notebook | 15.6 | Full HD 1920x1080 | Intel Core i5 7200U 2.5GHz | 8 | 256GB SSD | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 | 0 | 0 | |
| 3 | Apple | Ultrabook | 15.4 | IPS Panel Retina Display 2880x1800 | Intel Core i7 2.7GHz | 16 | 512GB SSD | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 | 0 | 1 | |
| 4 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 3.1GHz | 8 | 256GB SSD | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 | 0 | 1 | |



```
In [42]: 1 df['X_res'] = df['X_res'].astype('int')
2 df['Y_res'] = df['Y_res'].astype('int')
```

```
In [43]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1303 entries, 0 to 1302
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Company          1303 non-null    object  
 1   TypeName         1303 non-null    object  
 2   Inches           1303 non-null    float64 
 3   ScreenResolution 1303 non-null    object  
 4   Cpu              1303 non-null    object  
 5   Ram              1303 non-null    int32   
 6   Memory           1303 non-null    object  
 7   Gpu              1303 non-null    object  
 8   OpSys            1303 non-null    object  
 9   Weight            1303 non-null    float32 
 10  Price             1303 non-null    float64 
 11  Touchscreen      1303 non-null    int64   
 12  Ips              1303 non-null    int64   
 13  X_res            1303 non-null    int32   
 14  Y_res            1303 non-null    int32   
dtypes: float32(1), float64(2), int32(3), int64(2), object(7)
memory usage: 132.5+ KB
```

```
In [44]: 1 df.corr()['Price']
```

```
C:\Users\sanket parekh\AppData\Local\Temp\ipykernel_35592\815546952.py:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.
df.corr()['Price']
```

```
Out[44]: Inches      0.068197
Ram          0.743007
Weight        0.210370
Price         1.000000
Touchscreen   0.191226
Ips          0.252208
X_res         0.556529
Y_res         0.552809
Name: Price, dtype: float64
```

```
In [45]: 1 df['ppi'] = (((df['X_res']**2) + (df['Y_res']**2))**0.5/df['Inches']).astype('float')
2 df.corr()['Price']
```

C:\Users\sanket parekh\AppData\Local\Temp\ipykernel_35592\4248563634.py:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.
df.corr()['Price']

```
Out[45]: Inches      0.068197
Ram          0.743007
Weight       0.210370
Price        1.000000
Touchscreen  0.191226
Ips          0.252208
X_res        0.556529
Y_res        0.552809
ppi          0.473487
Name: Price, dtype: float64
```

In [46]:

```
1 df.drop(columns=['ScreenResolution'], inplace=True)
2 df.head(5)
```

Out[46]:

| | Company | TypeName | Inches | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscreen | Ips | X_res | Y_res | |
|---|---------|-----------|--------|----------------------------|-----|---------------------|------------------------------|-------|--------|-------------|-------------|-----|-------|-------|-------|
| 0 | Apple | Ultrabook | 13.3 | Intel Core i5 2.3GHz | 8 | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | 0 | 1 | 2560 | 1600 | 226.9 |
| 1 | Apple | Ultrabook | 13.3 | Intel Core i5 1.8GHz | 8 | 128GB Flash Storage | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 | 0 | 0 | 1440 | 900 | 127.6 |
| 2 | HP | Notebook | 15.6 | Intel Core i5 7200U 2.5GHz | 8 | 256GB SSD | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 | 0 | 0 | 1920 | 1080 | 141.2 |
| 3 | Apple | Ultrabook | 15.4 | Intel Core i7 2.7GHz | 16 | 512GB SSD | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 | 0 | 1 | 2880 | 1800 | 220.5 |
| 4 | Apple | Ultrabook | 13.3 | Intel Core i5 3.1GHz | 8 | 256GB SSD | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 | 0 | 1 | 2560 | 1600 | 226.9 |



```
In [47]: 1 df.drop(columns=['Inches','X_res','Y_res'],inplace=True)
2 df.head()
```

Out[47]:

| | Company | TypeName | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscreen | Ips | ppi |
|---|---------|-----------|----------------------------|-----|---------------------|------------------------------|-------|--------|-------------|-------------|-----|------------|
| 0 | Apple | Ultrabook | Intel Core i5 2.3GHz | 8 | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | 0 | 1 | 226.983005 |
| 1 | Apple | Ultrabook | Intel Core i5 1.8GHz | 8 | 128GB Flash Storage | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 | 0 | 0 | 127.677940 |
| 2 | HP | Notebook | Intel Core i5 7200U 2.5GHz | 8 | 256GB SSD | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 | 0 | 0 | 141.211998 |
| 3 | Apple | Ultrabook | Intel Core i7 2.7GHz | 16 | 512GB SSD | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 | 0 | 1 | 220.534624 |
| 4 | Apple | Ultrabook | Intel Core i5 3.1GHz | 8 | 256GB SSD | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 | 0 | 1 | 226.983005 |

```
In [48]: 1 df['Cpu'].value_counts()
```

Out[48]:

| | |
|--------------------------------------|-----|
| Intel Core i5 7200U 2.5GHz | 190 |
| Intel Core i7 7700HQ 2.8GHz | 146 |
| Intel Core i7 7500U 2.7GHz | 134 |
| Intel Core i7 8550U 1.8GHz | 73 |
| Intel Core i5 8250U 1.6GHz | 72 |
| ... | |
| Intel Core M M3-6Y30 0.9GHz | 1 |
| AMD A9-Series 9420 2.9GHz | 1 |
| Intel Core i3 6006U 2.2GHz | 1 |
| AMD A6-Series 7310 2GHz | 1 |
| Intel Xeon E3-1535M v6 3.1GHz | 1 |
| Name: Cpu, Length: 118, dtype: int64 | |

```
In [49]: 1 df['Cpu Name'] = df['Cpu'].apply(lambda x: " ".join(x.split()[0:3]))
2 df.head()
```

Out[49]:

| | Company | TypeName | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cpu Name |
|---|---------|-----------|----------------------------|-----|---------------------|------------------------------|-------|--------|-------------|-------------|-----|------------|---------------|
| 0 | Apple | Ultrabook | Intel Core i5 2.3GHz | 8 | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | 0 | 1 | 226.983005 | Intel Core i5 |
| 1 | Apple | Ultrabook | Intel Core i5 1.8GHz | 8 | 128GB Flash Storage | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 | 0 | 0 | 127.677940 | Intel Core i5 |
| 2 | HP | Notebook | Intel Core i5 7200U 2.5GHz | 8 | 256GB SSD | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 | 0 | 0 | 141.211998 | Intel Core i5 |
| 3 | Apple | Ultrabook | Intel Core i7 2.7GHz | 16 | 512GB SSD | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 | 0 | 1 | 220.534624 | Intel Core i7 |
| 4 | Apple | Ultrabook | Intel Core i5 3.1GHz | 8 | 256GB SSD | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 | 0 | 1 | 226.983005 | Intel Core i5 |

```
In [50]: 1 def fetch_processor(text):
2     if text == 'Intel Core i7' or text == 'Intel Core i5' or text == 'Intel Core i3':
3         return text
4     else:
5         if text.split()[0] == 'Intel':
6             return 'Other Intel Processor'
7         else:
8             return 'AMD Processor'
```

In [51]:

```

1 df['Cpu brand'] = df['Cpu Name'].apply(fetch_processor)
2 df.head()

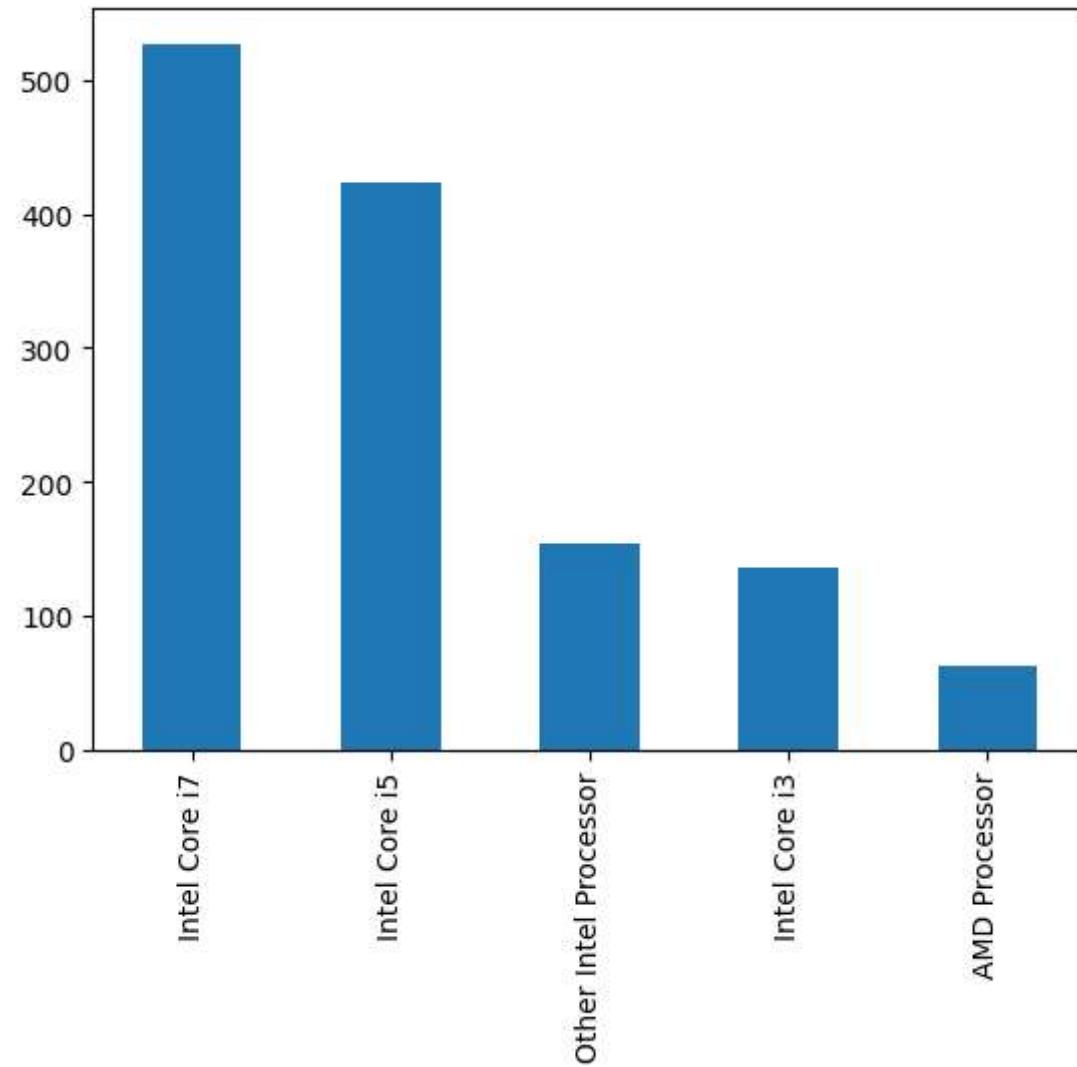
```

Out[51]:

| | Company | TypeName | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cpu Name | Cpu brand |
|---|---------|-----------|----------------------------|-----|---------------------|------------------------------|-------|--------|-------------|-------------|-----|------------|---------------|---------------|
| 0 | Apple | Ultrabook | Intel Core i5 2.3GHz | 8 | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | 0 | 1 | 226.983005 | Intel Core i5 | Intel Core i5 |
| 1 | Apple | Ultrabook | Intel Core i5 1.8GHz | 8 | 128GB Flash Storage | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 | 0 | 0 | 127.677940 | Intel Core i5 | Intel Core i5 |
| 2 | HP | Notebook | Intel Core i5 7200U 2.5GHz | 8 | 256GB SSD | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 | 0 | 0 | 141.211998 | Intel Core i5 | Intel Core i5 |
| 3 | Apple | Ultrabook | Intel Core i7 2.7GHz | 16 | 512GB SSD | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 | 0 | 1 | 220.534624 | Intel Core i7 | Intel Core i7 |
| 4 | Apple | Ultrabook | Intel Core i5 3.1GHz | 8 | 256GB SSD | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 | 0 | 1 | 226.983005 | Intel Core i5 | Intel Core i5 |

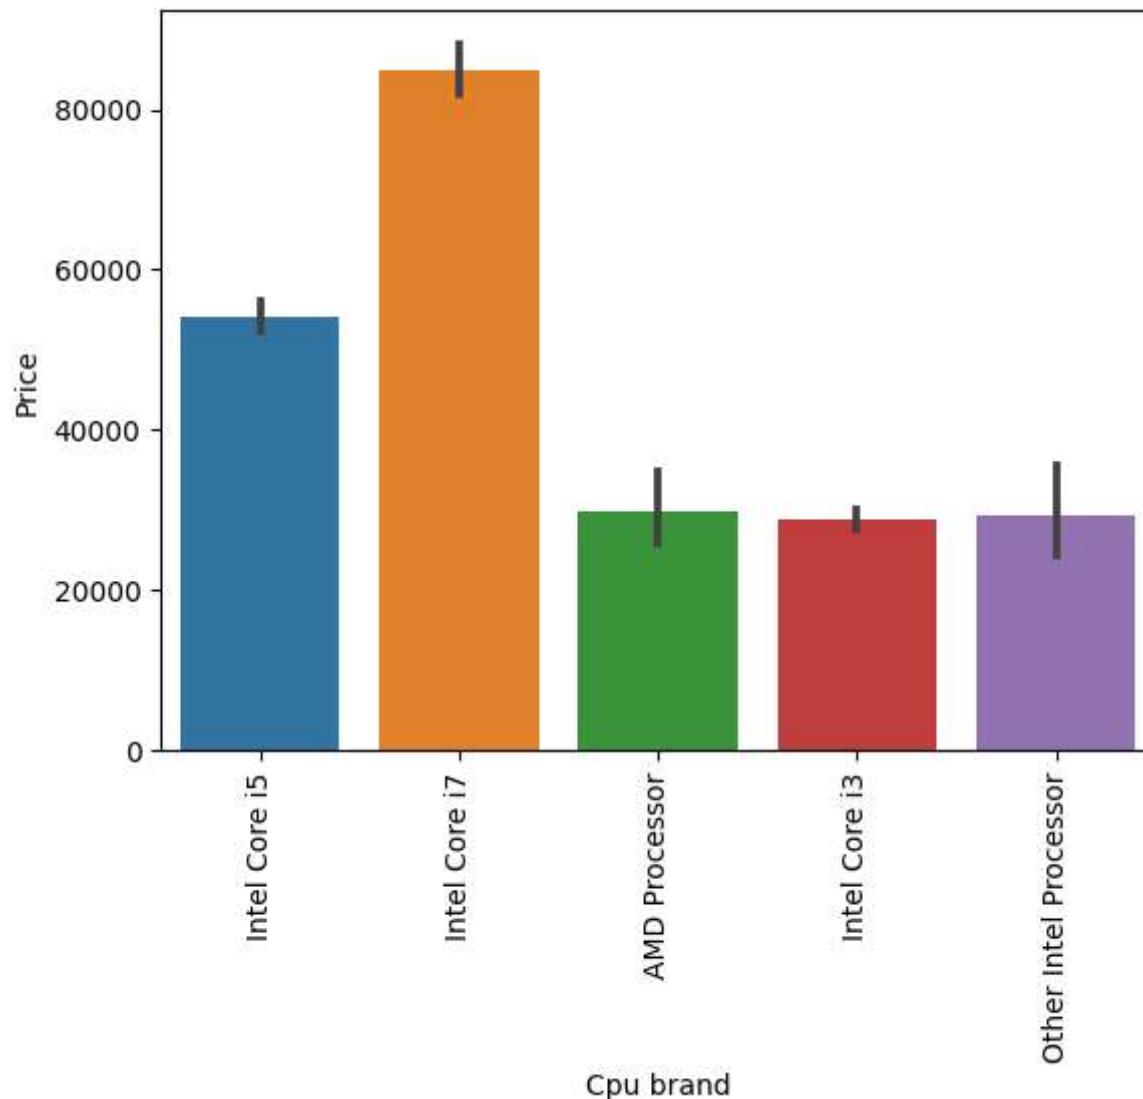
```
In [52]: 1 df['cpu brand'].value_counts().plot(kind='bar')
```

```
Out[52]: <AxesSubplot: >
```



In [53]:

```
1 sns.barplot(x=df['Cpu brand'],y=df['Price'])
2 plt.xticks(rotation='vertical')
3 plt.show()
```



```
In [54]: 1 df.drop(columns=['Cpu', 'Cpu Name'], inplace=True)
```

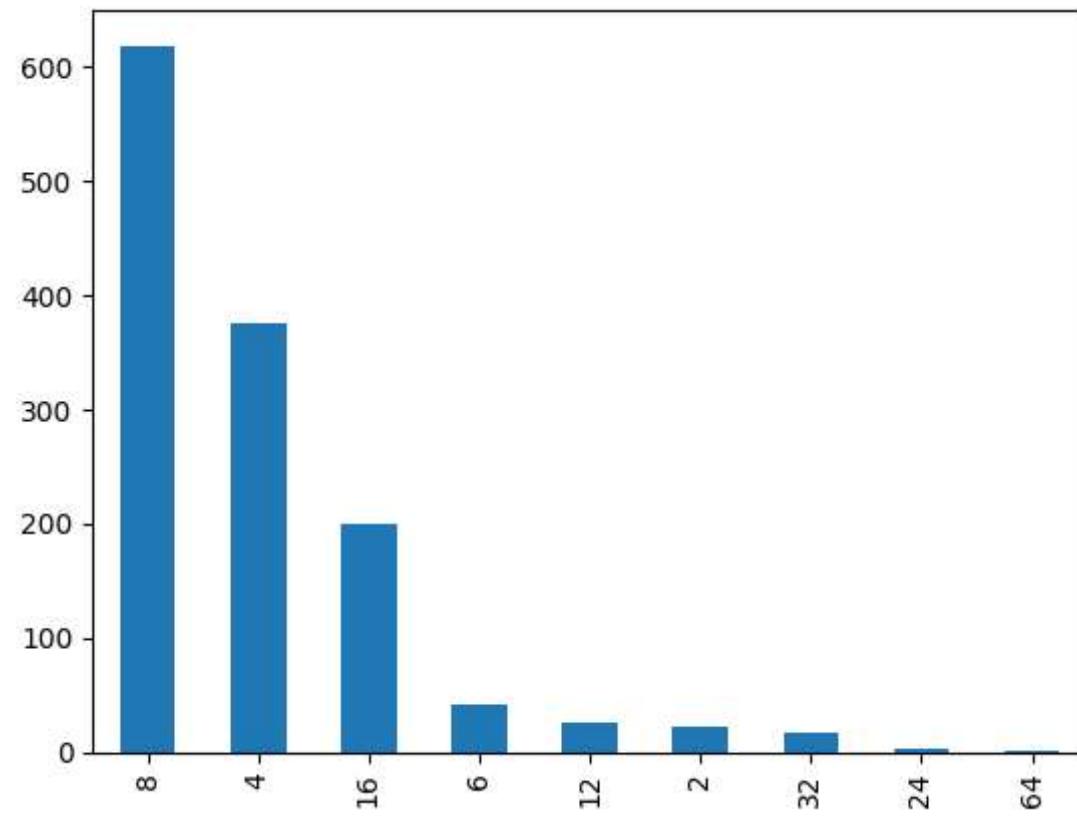
```
In [55]: 1 df.head()
```

Out[55]:

| | Company | TypeName | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cpu brand |
|---|---------|-----------|-----|---------------------|------------------------------|-------|--------|-------------|-------------|-----|------------|---------------|
| 0 | Apple | Ultrabook | 8 | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | 0 | 1 | 226.983005 | Intel Core i5 |
| 1 | Apple | Ultrabook | 8 | 128GB Flash Storage | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 | 0 | 0 | 127.677940 | Intel Core i5 |
| 2 | HP | Notebook | 8 | 256GB SSD | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 | 0 | 0 | 141.211998 | Intel Core i5 |
| 3 | Apple | Ultrabook | 16 | 512GB SSD | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 | 0 | 1 | 220.534624 | Intel Core i7 |
| 4 | Apple | Ultrabook | 8 | 256GB SSD | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 | 0 | 1 | 226.983005 | Intel Core i5 |

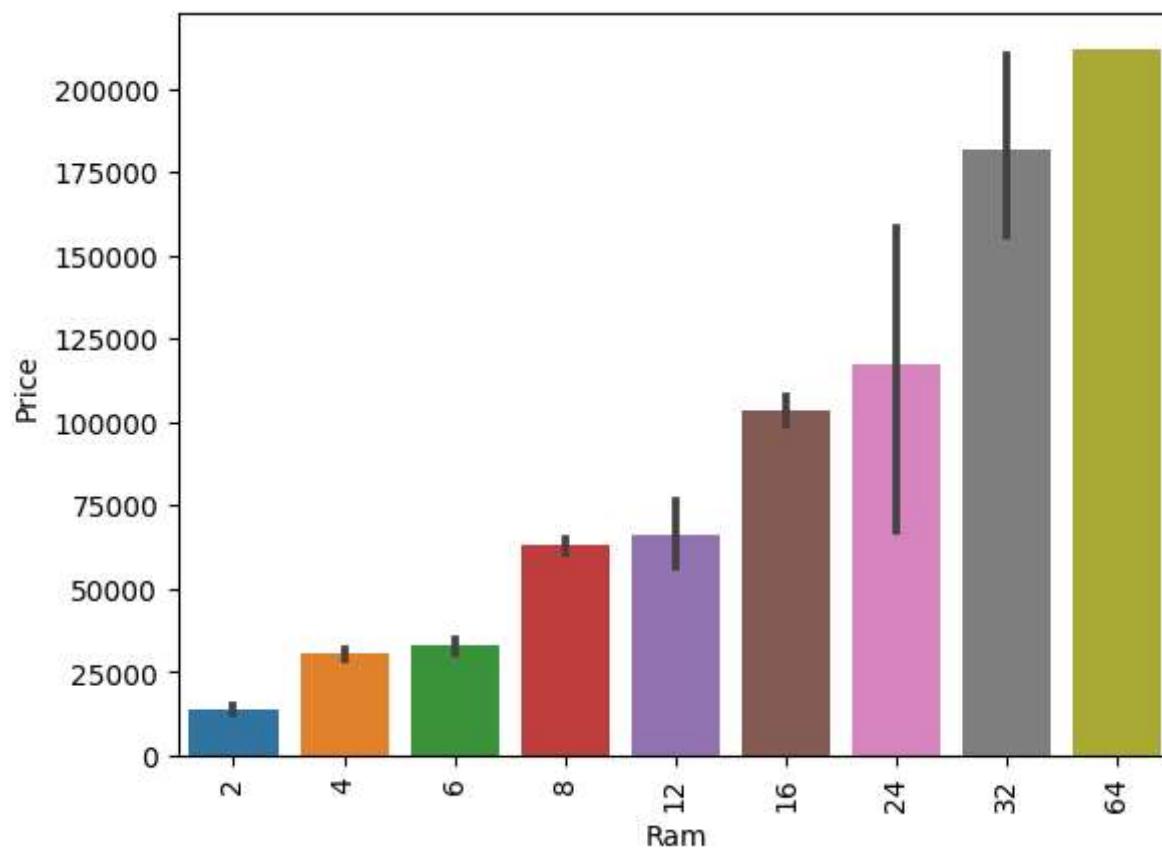
```
In [56]: 1 df['Ram'].value_counts().plot(kind='bar')
```

```
Out[56]: <AxesSubplot: >
```



In [57]:

```
1 sns.barplot(x=df['Ram'],y=df['Price'])
2 plt.xticks(rotation='vertical')
3 plt.show()
```



```
In [58]: 1 df['Memory'].value_counts()
```

```
Out[58]: 256GB SSD           412
1TB HDD              223
500GB HDD             132
512GB SSD             118
128GB SSD + 1TB HDD    94
128GB SSD              76
256GB SSD + 1TB HDD    73
32GB Flash Storage     38
2TB HDD                16
64GB Flash Storage      15
512GB SSD + 1TB HDD    14
1TB SSD                14
256GB SSD + 2TB HDD     10
1.0TB Hybrid            9
256GB Flash Storage      8
16GB Flash Storage       7
32GB SSD                6
180GB SSD               5
128GB Flash Storage      4
512GB SSD + 2TB HDD      3
16GB SSD                3
512GB Flash Storage      2
1TB SSD + 1TB HDD        2
256GB SSD + 500GB HDD      2
128GB SSD + 2TB HDD        2
256GB SSD + 256GB SSD      2
512GB SSD + 256GB SSD      1
512GB SSD + 512GB SSD      1
64GB Flash Storage + 1TB HDD  1
1TB HDD + 1TB HDD          1
32GB HDD                1
64GB SSD                1
128GB HDD                1
240GB SSD                1
8GB SSD                 1
508GB Hybrid              1
1.0TB HDD                1
512GB SSD + 1.0TB Hybrid    1
256GB SSD + 1.0TB Hybrid    1
Name: Memory, dtype: int64
```

In [59]:

```
1 df['Memory'] = df['Memory'].astype(str).replace('.0', '', regex=True)
2 df["Memory"] = df["Memory"].str.replace('GB', '')
3 df["Memory"] = df["Memory"].str.replace('TB', '000')
4 new = df["Memory"].str.split("+", n = 1, expand = True)
5
6 df["first"] = new[0]
7 df["first"] = df["first"].str.strip()
8
9 df["second"] = new[1]
10
11 df["Layer1HDD"] = df["first"].apply(lambda x: 1 if "HDD" in x else 0)
12 df["Layer1SSD"] = df["first"].apply(lambda x: 1 if "SSD" in x else 0)
13 df["Layer1Hybrid"] = df["first"].apply(lambda x: 1 if "Hybrid" in x else 0)
14 df["Layer1Flash_Storage"] = df["first"].apply(lambda x: 1 if "Flash Storage" in x else 0)
15
16 df['first'] = df['first'].str.replace(r'\D', '')
17
18 df["second"].fillna("0", inplace = True)
19
20 df["Layer2HDD"] = df["second"].apply(lambda x: 1 if "HDD" in x else 0)
21 df["Layer2SSD"] = df["second"].apply(lambda x: 1 if "SSD" in x else 0)
22 df["Layer2Hybrid"] = df["second"].apply(lambda x: 1 if "Hybrid" in x else 0)
23 df["Layer2Flash_Storage"] = df["second"].apply(lambda x: 1 if "Flash Storage" in x else 0)
24
25 df['second'] = df['second'].str.replace(r'\D', '')
26
27 df["first"] = df["first"].astype(int)
28 df["second"] = df["second"].astype(int)
29
30 df["HDD"]=(df["first"]*df["Layer1HDD"]+df["second"]*df["Layer2HDD"])
31 df["SSD"]=(df["first"]*df["Layer1SSD"]+df["second"]*df["Layer2SSD"])
32 df["Hybrid"]=(df["first"]*df["Layer1Hybrid"]+df["second"]*df["Layer2Hybrid"])
33 df["Flash_Storage"]=(df["first"]*df["Layer1Flash_Storage"]+df["second"]*df["Layer2Flash_Storage"])
34
35 df.drop(columns=['first', 'second', 'Layer1HDD', 'Layer1SSD', 'Layer1Hybrid',
36             'Layer1Flash_Storage', 'Layer2HDD', 'Layer2SSD', 'Layer2Hybrid',
37             'Layer2Flash_Storage'], inplace=True)
```

```
C:\Users\sanket parekh\AppData\Local\Temp\ipykernel_35592\4023190604.py:16: FutureWarning: The default value  
of regex will change from True to False in a future version.  
    df['first'] = df['first'].str.replace(r'\D', '')  
C:\Users\sanket parekh\AppData\Local\Temp\ipykernel_35592\4023190604.py:25: FutureWarning: The default value  
of regex will change from True to False in a future version.  
    df['second'] = df['second'].str.replace(r'\D', '')
```

In [60]: 1 df.sample(5)

Out[60]:

| | Company | TypeName | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cpu brand | HDD | SSD | H |
|-----|---------|--------------------|-----|--------------------|-------------------------|------------|--------|------------|-------------|-----|------------|---------------|------|-----|---|
| 919 | MSI | Gaming | 8 | 128 SSD + 1000 HDD | Nvidia GeForce GTX 1050 | Windows 10 | 2.20 | 59668.8048 | 0 | 0 | 141.211998 | Intel Core i5 | 1000 | 128 | |
| 880 | HP | 2 in 1 Convertible | 4 | 256 SSD | Intel HD Graphics 620 | Windows 10 | 1.28 | 90576.0000 | 1 | 0 | 165.632118 | Intel Core i5 | 0 | 256 | |
| 902 | Dell | 2 in 1 Convertible | 16 | 256 SSD | Intel HD Graphics 615 | Windows 10 | 1.22 | 87858.7200 | 1 | 0 | 165.632118 | Intel Core i7 | 0 | 256 | |
| 624 | HP | Notebook | 4 | 500 HDD | Intel HD Graphics 520 | Windows 7 | 1.88 | 49656.9600 | 0 | 0 | 100.454670 | Intel Core i5 | 500 | 0 | |
| 55 | Dell | Notebook | 8 | 256 SSD | AMD Radeon 520 | Windows 10 | 2.13 | 40908.3840 | 0 | 0 | 141.211998 | Intel Core i7 | 0 | 256 | |



```
In [61]: 1 df.drop(columns=['Memory'], inplace=True)
2 df.head()
```

Out[61]:

| | Company | TypeName | Ram | Gpu | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cpu brand | HDD | SSD | Hybrid | Flash_ |
|---|---------|-----------|-----|------------------------------|-------|--------|-------------|-------------|-----|------------|---------------|-----|-----|--------|--------|
| 0 | Apple | Ultrabook | 8 | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | | 0 1 | 226.983005 | Intel Core i5 | 0 | 128 | 0 | |
| 1 | Apple | Ultrabook | 8 | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 | | 0 0 | 127.677940 | Intel Core i5 | 0 | 0 | 0 | |
| 2 | HP | Notebook | 8 | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 | | 0 0 | 141.211998 | Intel Core i5 | 0 | 256 | 0 | |
| 3 | Apple | Ultrabook | 16 | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 | | 0 1 | 220.534624 | Intel Core i7 | 0 | 512 | 0 | |
| 4 | Apple | Ultrabook | 8 | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 | | 0 1 | 226.983005 | Intel Core i5 | 0 | 256 | 0 | |



```
In [62]: 1 df.drop(columns=['Hybrid','Flash_Storage'],inplace=True)
2 df.head()
```

Out[62]:

| | Company | TypeName | Ram | Gpu | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cpu brand | HDD | SSD |
|---|---------|-----------|-----|------------------------------|-------|--------|-------------|-------------|-----|------------|---------------|-----|-----|
| 0 | Apple | Ultrabook | 8 | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | | 0 1 | 226.983005 | Intel Core i5 | 0 | 128 |
| 1 | Apple | Ultrabook | 8 | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 | | 0 0 | 127.677940 | Intel Core i5 | 0 | 0 |
| 2 | HP | Notebook | 8 | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 | | 0 0 | 141.211998 | Intel Core i5 | 0 | 256 |
| 3 | Apple | Ultrabook | 16 | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 | | 0 1 | 220.534624 | Intel Core i7 | 0 | 512 |
| 4 | Apple | Ultrabook | 8 | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 | | 0 1 | 226.983005 | Intel Core i5 | 0 | 256 |

```
In [63]: 1 df['Gpu'].value_counts()
```

Out[63]:

| | |
|-------------------------|-----|
| Intel HD Graphics 620 | 281 |
| Intel HD Graphics 520 | 185 |
| Intel UHD Graphics 620 | 68 |
| Nvidia GeForce GTX 1050 | 66 |
| Nvidia GeForce GTX 1060 | 48 |
| ... | |
| AMD Radeon R5 520 | 1 |
| AMD Radeon R7 | 1 |
| Intel HD Graphics 540 | 1 |
| AMD Radeon 540 | 1 |
| ARM Mali T860 MP4 | 1 |

Name: Gpu, Length: 110, dtype: int64

```
In [64]: 1 df['Gpu brand'] = df['Gpu'].apply(lambda x:x.split()[0])
2 df.head()
```

Out[64]:

| | Company | TypeName | Ram | Gpu | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cpu brand | HDD | SSD | Gpu brand |
|---|---------|-----------|-----|------------------------------|-------|--------|-------------|-------------|-----|------------|---------------|-----|-----|-----------|
| 0 | Apple | Ultrabook | 8 | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | | 0 1 | 226.983005 | Intel Core i5 | 0 | 128 | Intel |
| 1 | Apple | Ultrabook | 8 | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 | | 0 0 | 127.677940 | Intel Core i5 | 0 | 0 | Intel |
| 2 | HP | Notebook | 8 | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 | | 0 0 | 141.211998 | Intel Core i5 | 0 | 256 | Intel |
| 3 | Apple | Ultrabook | 16 | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 | | 0 1 | 220.534624 | Intel Core i7 | 0 | 512 | AMD |
| 4 | Apple | Ultrabook | 8 | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 | | 0 1 | 226.983005 | Intel Core i5 | 0 | 256 | Intel |

```
In [65]: 1 df['Gpu brand'].value_counts()
```

Out[65]: Intel 722
Nvidia 400
AMD 180
ARM 1
Name: Gpu brand, dtype: int64

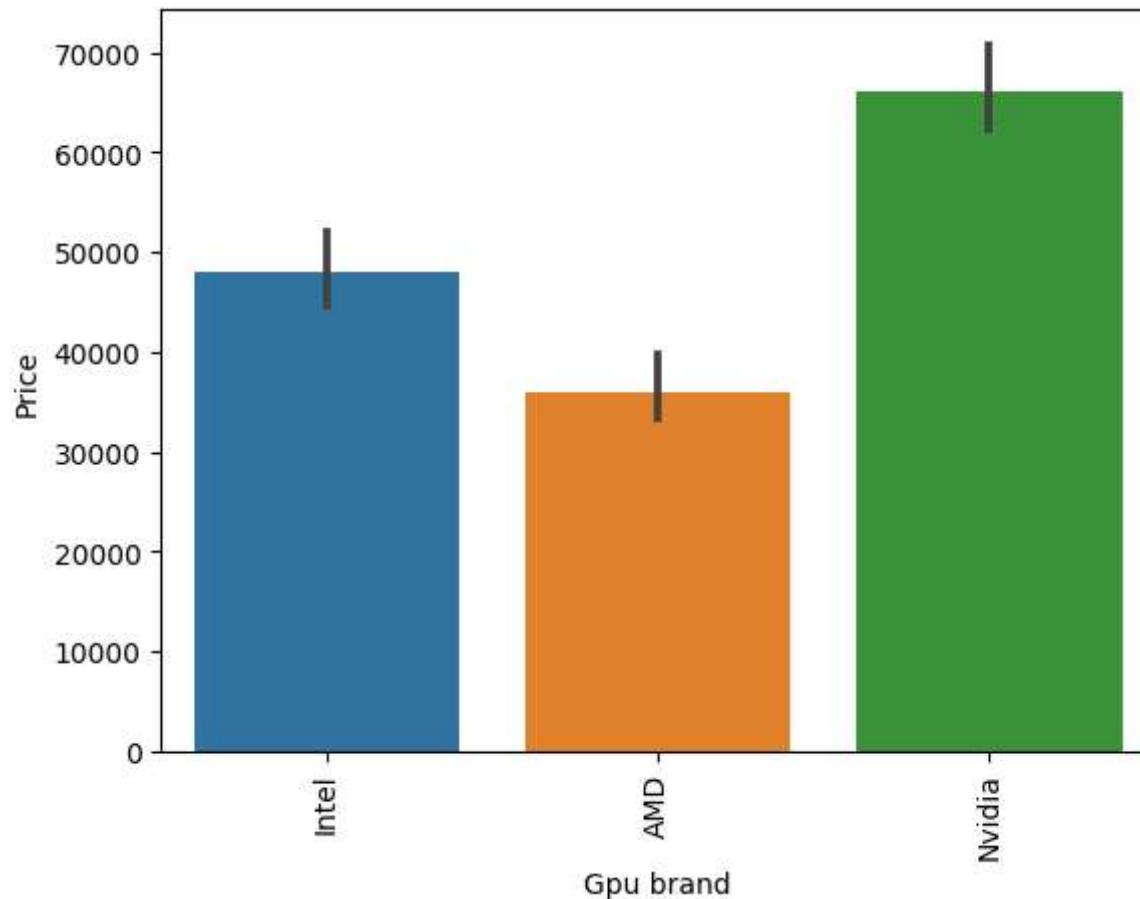
```
In [66]: 1 df = df[df['Gpu brand'] != 'ARM']
```

```
In [67]: 1 df['Gpu brand'].value_counts()
```

Out[67]: Intel 722
Nvidia 400
AMD 180
Name: Gpu brand, dtype: int64

In [68]:

```
1 sns.barplot(x=df['Gpu brand'],y=df['Price'],estimator=np.median)
2 plt.xticks(rotation='vertical')
3 plt.show()
```



In [69]:

```
1 df.drop(columns=['Gpu'],inplace=True)
```

C:\Users\sanket parekh\AppData\Local\Temp\ipykernel_35592\1111925144.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

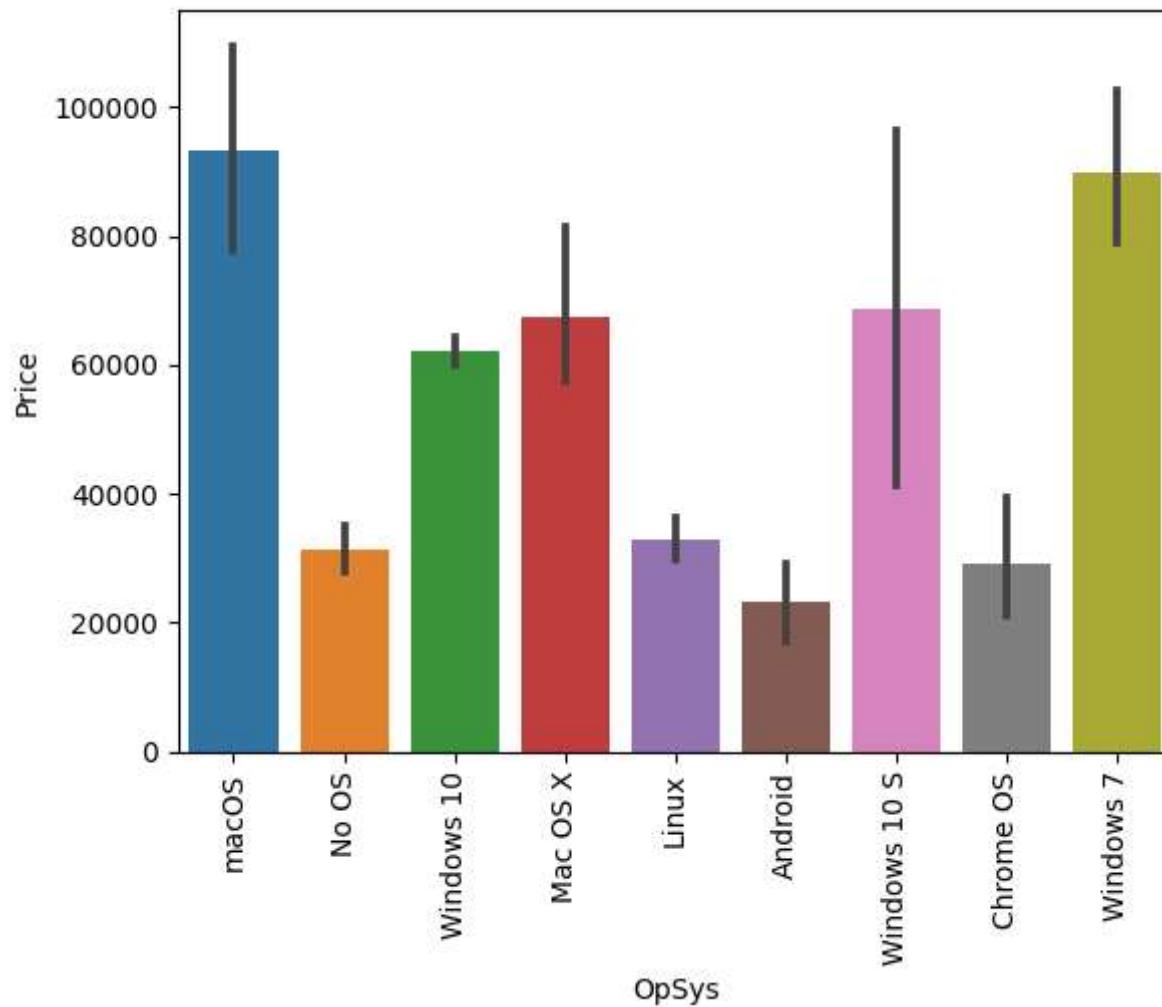
```
df.drop(columns=['Gpu'],inplace=True)
```

```
In [70]: 1 df['OpSys'].value_counts()
```

```
Out[70]: Windows 10      1072  
No OS                 66  
Linux                 62  
Windows 7              45  
Chrome OS              26  
macOS                  13  
Mac OS X                8  
Windows 10 S              8  
Android                  2  
Name: OpSys, dtype: int64
```

In [71]:

```
1 sns.barplot(x=df['OpSys'],y=df['Price'])
2 plt.xticks(rotation='vertical')
3 plt.show()
```



```
In [72]: 1 def cat_os(inp):
2     if inp == 'Windows 10' or inp == 'Windows 7' or inp == 'Windows 10 S':
3         return 'Windows'
4     elif inp == 'macOS' or inp == 'Mac OS X':
5         return 'Mac'
6     else:
7         return 'Others/No OS/Linux'
```

```
In [73]: 1 df['os'] = df['OpSys'].apply(cat_os)
2 df.head()
```

C:\Users\sanket parekh\AppData\Local\Temp\ipykernel_35592\246390668.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df['os'] = df['OpSys'].apply(cat_os)
```

Out[73]:

| | Company | TypeName | Ram | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cpu brand | HDD | SSD | Gpu brand | os | |
|---|---------|-----------|-----|-------|--------|-------------|-------------|-----|-----|------------|---------------|-----|-----------|-------|--------------------|
| 0 | Apple | Ultrabook | 8 | macOS | 1.37 | 71378.6832 | | 0 | 1 | 226.983005 | Intel Core i5 | 0 | 128 | Intel | Mac |
| 1 | Apple | Ultrabook | 8 | macOS | 1.34 | 47895.5232 | | 0 | 0 | 127.677940 | Intel Core i5 | 0 | 0 | Intel | Mac |
| 2 | HP | Notebook | 8 | No OS | 1.86 | 30636.0000 | | 0 | 0 | 141.211998 | Intel Core i5 | 0 | 256 | Intel | Others/No OS/Linux |
| 3 | Apple | Ultrabook | 16 | macOS | 1.83 | 135195.3360 | | 0 | 1 | 220.534624 | Intel Core i7 | 0 | 512 | AMD | Mac |
| 4 | Apple | Ultrabook | 8 | macOS | 1.37 | 96095.8080 | | 0 | 1 | 226.983005 | Intel Core i5 | 0 | 256 | Intel | Mac |

```
In [74]: 1 df.drop(columns=['OpSys'], inplace=True)
```

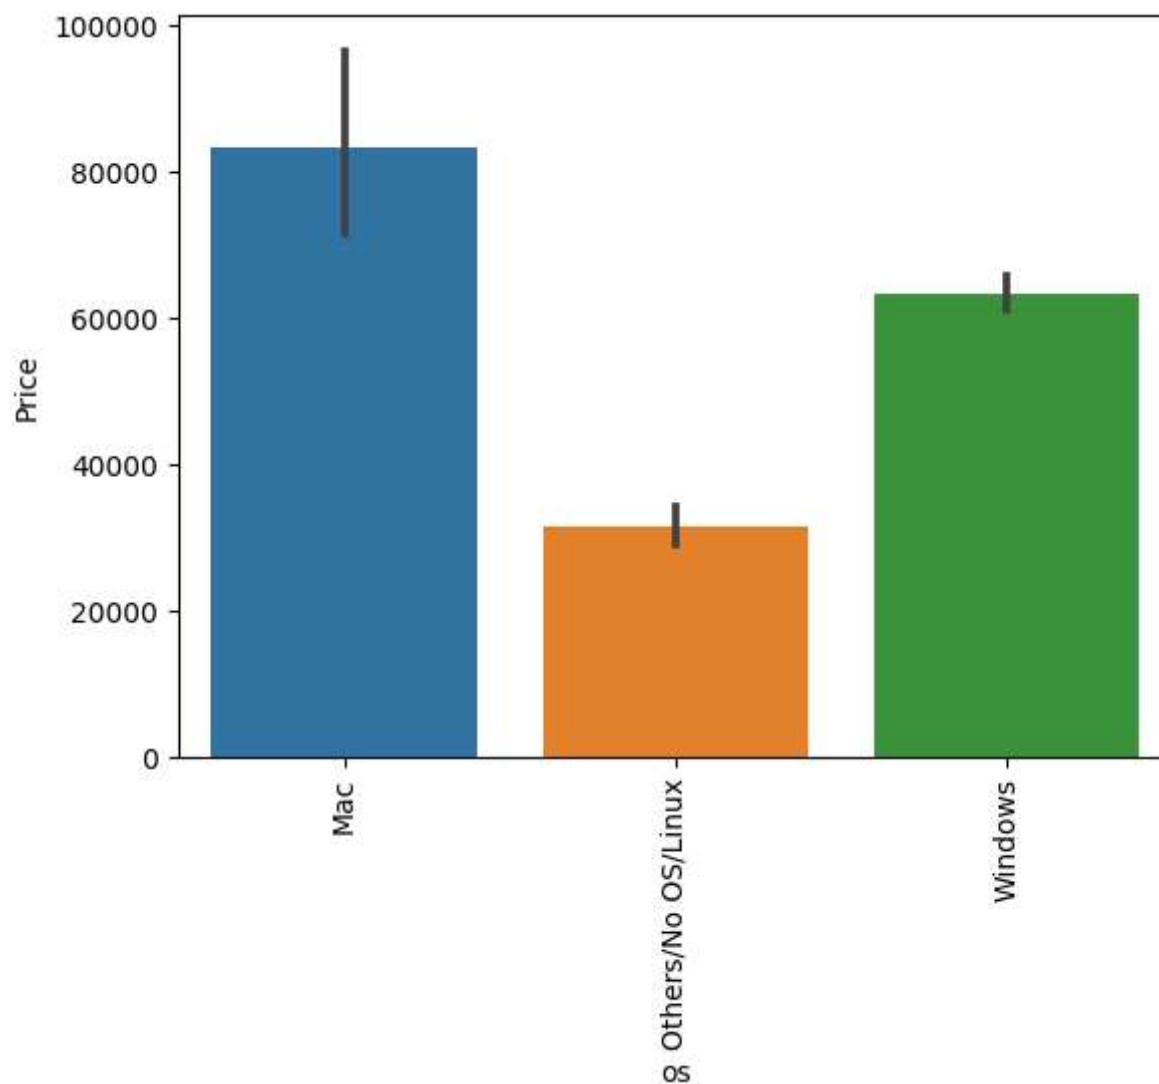
C:\Users\sanket parekh\AppData\Local\Temp\ipykernel_35592\3105339334.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df.drop(columns=['OpSys'], inplace=True)
```

In [75]:

```
1 sns.barplot(x=df['os'],y=df['Price'])
2 plt.xticks(rotation='vertical')
3 plt.show()
```



```
In [76]: 1 sns.distplot(df['Weight'])
```

C:\Users\sanket parekh\AppData\Local\Temp\ipykernel_35592\1125578356.py:1: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

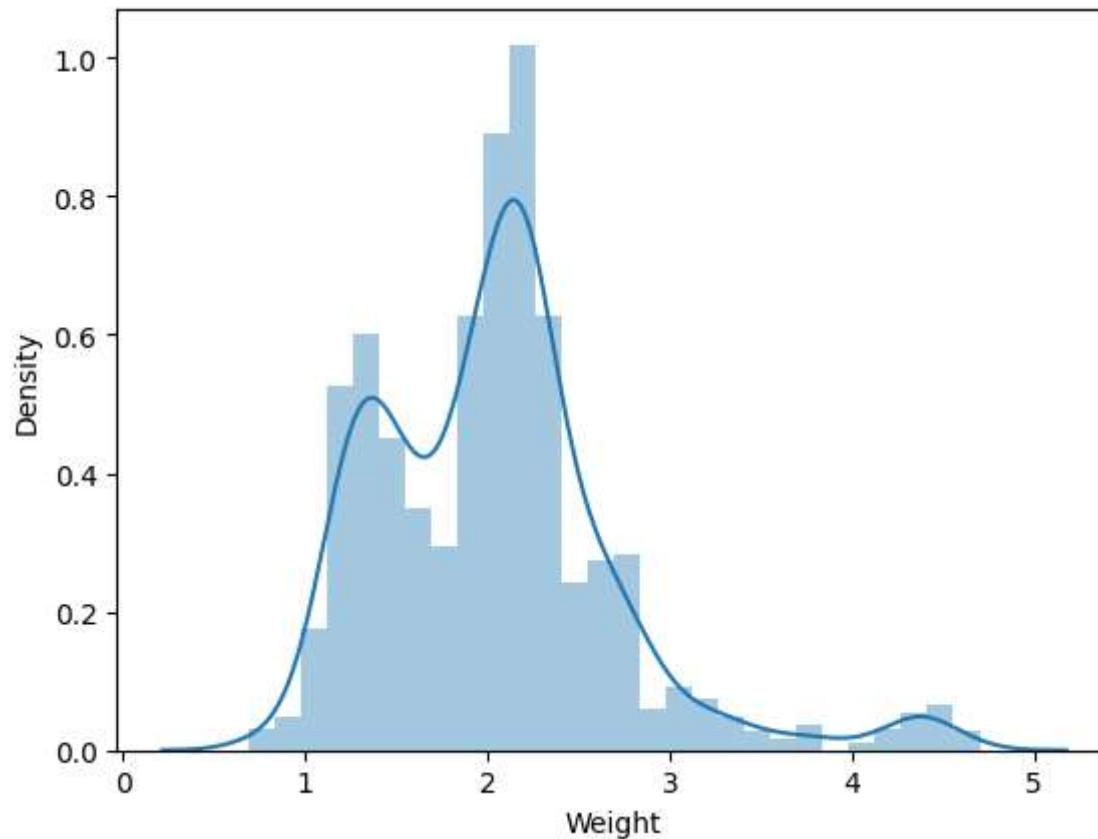
Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751> (<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>)

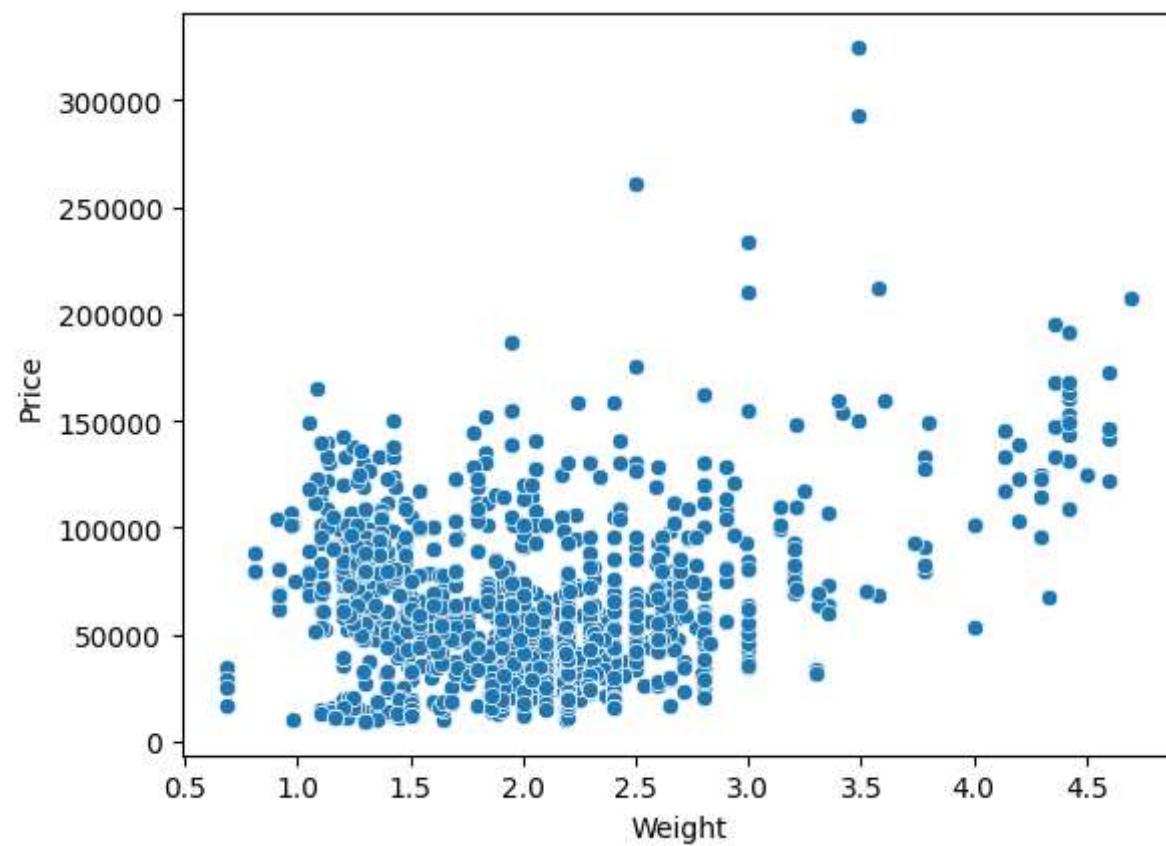
```
sns.distplot(df['Weight'])
```

Out[76]: <AxesSubplot: xlabel='Weight', ylabel='Density'>



```
In [77]: 1 sns.scatterplot(x=df['Weight'],y=df['Price'])
```

```
Out[77]: <AxesSubplot: xlabel='Weight', ylabel='Price'>
```

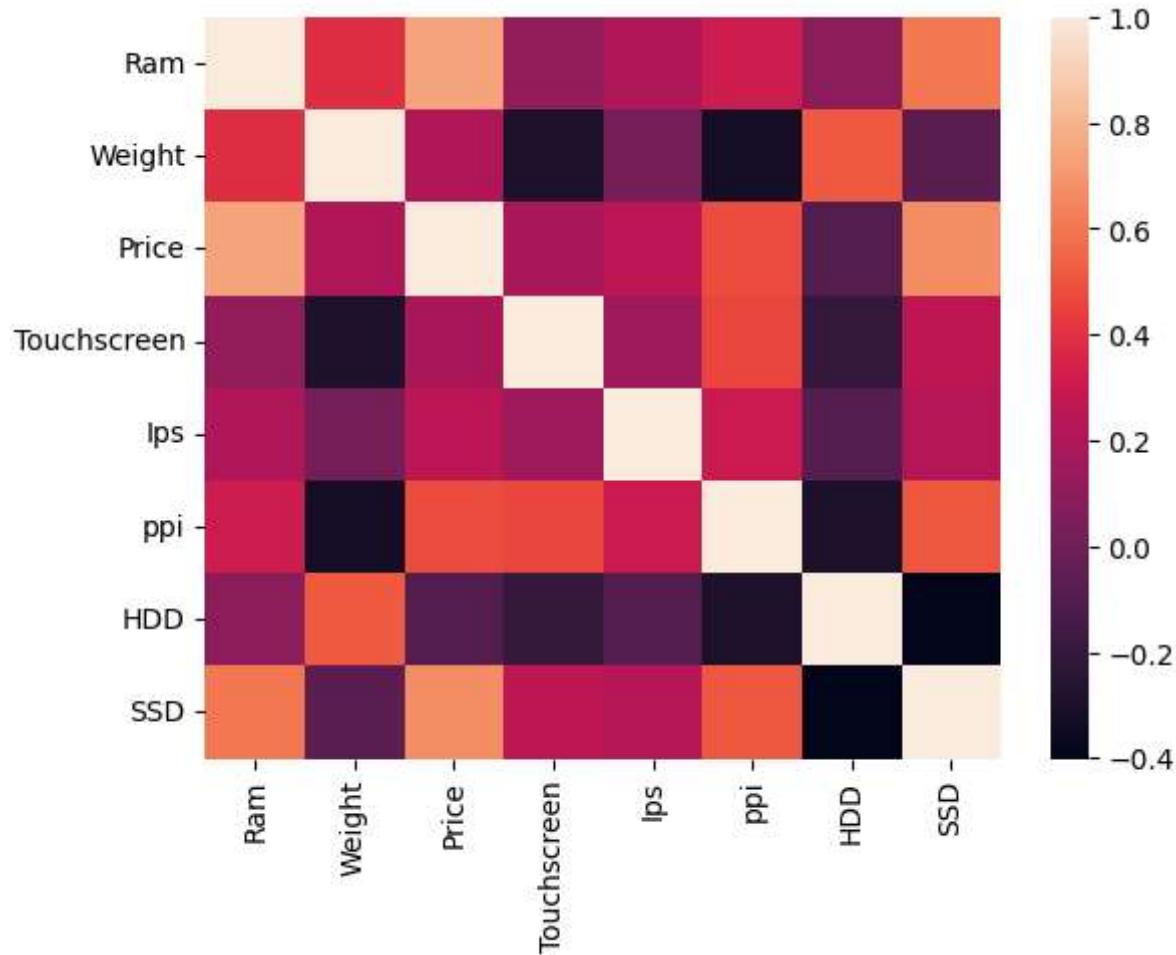


```
In [78]: 1 sns.heatmap(df.corr())  
          2
```

C:\Users\sanket parekh\AppData\Local\Temp\ipykernel_35592\1072140413.py:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
sns.heatmap(df.corr())
```

Out[78]: <AxesSubplot: >



```
In [79]: 1 sns.distplot(np.log(df['Price']))
```

C:\Users\sanket parekh\AppData\Local\Temp\ipykernel_35592\3556049916.py:1: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

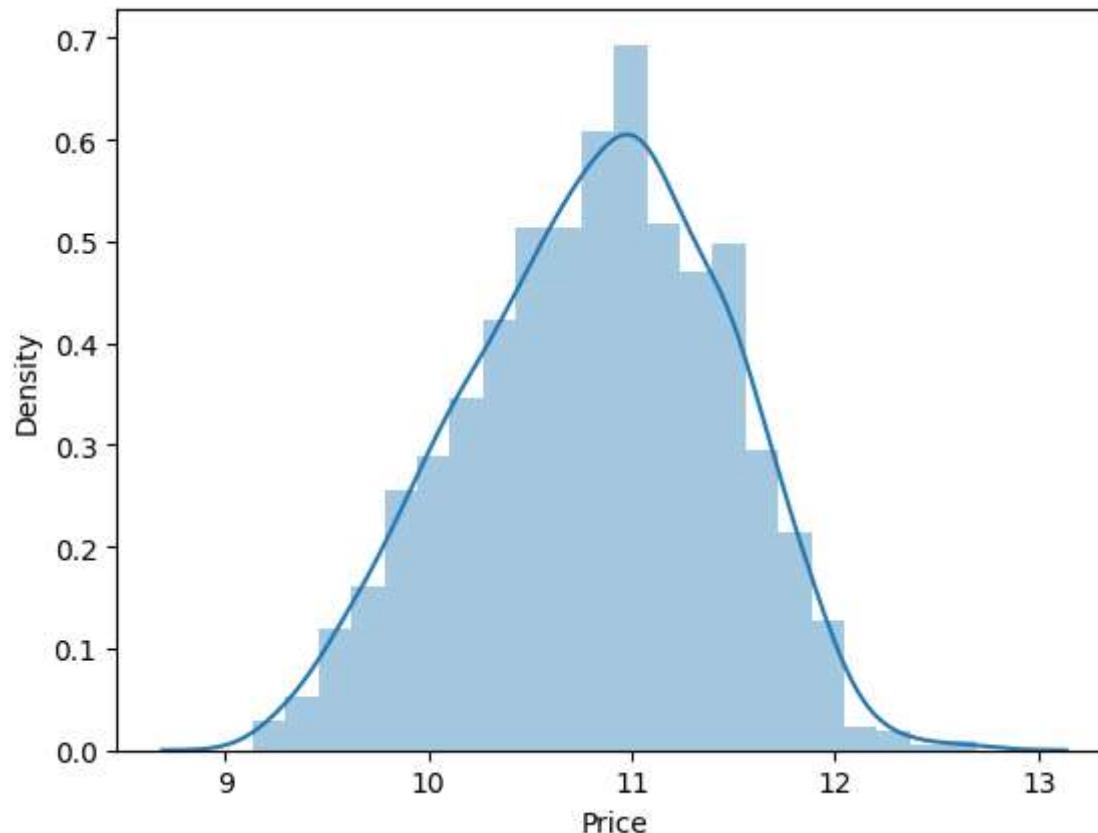
Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751> (<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>)

```
sns.distplot(np.log(df['Price']))
```

Out[79]: <AxesSubplot: xlabel='Price', ylabel='Density'>



```
In [80]: 1 X = df.drop(columns=['Price'])  
2 y = np.log(df['Price'])
```

```
In [82]: 1 X
```

Out[82]:

| | Company | TypeName | Ram | Weight | Touchscreen | Ips | ppi | Cpu brand | HDD | SSD | Gpu brand | os |
|------|---------|--------------------|-----|--------|-------------|-----|------------|-----------------------|------|-----|-----------|--------------------|
| 0 | Apple | Ultrabook | 8 | 1.37 | 0 | 1 | 226.983005 | Intel Core i5 | 0 | 128 | Intel | Mac |
| 1 | Apple | Ultrabook | 8 | 1.34 | 0 | 0 | 127.677940 | Intel Core i5 | 0 | 0 | Intel | Mac |
| 2 | HP | Notebook | 8 | 1.86 | 0 | 0 | 141.211998 | Intel Core i5 | 0 | 256 | Intel | Others/No OS/Linux |
| 3 | Apple | Ultrabook | 16 | 1.83 | 0 | 1 | 220.534624 | Intel Core i7 | 0 | 512 | AMD | Mac |
| 4 | Apple | Ultrabook | 8 | 1.37 | 0 | 1 | 226.983005 | Intel Core i5 | 0 | 256 | Intel | Mac |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1298 | Lenovo | 2 in 1 Convertible | 4 | 1.80 | 1 | 1 | 157.350512 | Intel Core i7 | 0 | 128 | Intel | Windows |
| 1299 | Lenovo | 2 in 1 Convertible | 16 | 1.30 | 1 | 1 | 276.053530 | Intel Core i7 | 0 | 512 | Intel | Windows |
| 1300 | Lenovo | Notebook | 2 | 1.50 | 0 | 0 | 111.935204 | Other Intel Processor | 0 | 0 | Intel | Windows |
| 1301 | HP | Notebook | 6 | 2.19 | 0 | 0 | 100.454670 | Intel Core i7 | 1000 | 0 | AMD | Windows |
| 1302 | Asus | Notebook | 4 | 2.20 | 0 | 0 | 100.454670 | Other Intel Processor | 500 | 0 | Intel | Windows |

1302 rows × 12 columns

```
In [83]: 1 y
```

```
Out[83]: 0      11.175755
1      10.776777
2      10.329931
3      11.814476
4      11.473101
...
1298    10.433899
1299    11.288115
1300    9.409283
1301    10.614129
1302    9.886358
Name: Price, Length: 1302, dtype: float64
```

```
In [84]: 1 from sklearn.model_selection import train_test_split
2 X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.15,random_state=2)
```

```
In [85]: 1 X_train
```

Out[85]:

| | Company | Type Name | Ram | Weight | Touchscreen | Lps | ppi | Cpu brand | HDD | SSD | Gpu brand | os |
|------|---------|--------------------|-----|--------|-------------|-----|------------|-----------------------|------|-----|-----------|--------------------|
| 183 | Toshiba | Notebook | 8 | 2.00 | 0 | 0 | 100.454670 | Intel Core i5 | 0 | 128 | Intel | Windows |
| 1141 | MSI | Gaming | 8 | 2.40 | 0 | 0 | 141.211998 | Intel Core i7 | 1000 | 128 | Nvidia | Windows |
| 1049 | Asus | Netbook | 4 | 1.20 | 0 | 0 | 135.094211 | Other Intel Processor | 0 | 0 | Intel | Others/No OS/Linux |
| 1020 | Dell | 2 in 1 Convertible | 4 | 2.08 | 1 | 1 | 141.211998 | Intel Core i3 | 1000 | 0 | Intel | Windows |
| 878 | Dell | Notebook | 4 | 2.18 | 0 | 0 | 141.211998 | Intel Core i5 | 1000 | 128 | Nvidia | Windows |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 466 | Acer | Notebook | 4 | 2.20 | 0 | 0 | 100.454670 | Intel Core i3 | 500 | 0 | Nvidia | Windows |
| 299 | Asus | Ultrabook | 16 | 1.63 | 0 | 0 | 141.211998 | Intel Core i7 | 0 | 512 | Nvidia | Windows |
| 493 | Acer | Notebook | 8 | 2.20 | 0 | 0 | 100.454670 | AMD Processor | 1000 | 0 | AMD | Windows |
| 527 | Lenovo | Notebook | 8 | 2.20 | 0 | 0 | 100.454670 | Intel Core i3 | 2000 | 0 | Nvidia | Others/No OS/Linux |
| 1193 | Apple | Ultrabook | 8 | 0.92 | 0 | 1 | 226.415547 | Other Intel Processor | 0 | 0 | Intel | Mac |

1106 rows × 12 columns

```
In [86]: 1 from sklearn.compose import ColumnTransformer
2 from sklearn.pipeline import Pipeline
3 from sklearn.preprocessing import OneHotEncoder
4 from sklearn.metrics import r2_score, mean_absolute_error
```

In [87]:

```
1 from sklearn.linear_model import LinearRegression,Ridge,Lasso
2 from sklearn.neighbors import KNeighborsRegressor
3 from sklearn.tree import DecisionTreeRegressor
4 from sklearn.ensemble import RandomForestRegressor,GradientBoostingRegressor,AdaBoostRegressor,ExtraTrees
5 from sklearn.svm import SVR
6 from xgboost import XGBRegressor
```

Linear Regression

In [88]:

```
1 step1 = ColumnTransformer(transformers=[
2     ('col_tnf',OneHotEncoder(sparse=False,drop='first'),[0,1,7,10,11])
3 ],remainder='passthrough')
4
5 step2 = LinearRegression()
6
7 pipe = Pipeline([
8     ('step1',step1),
9     ('step2',step2)
10])
11
12 pipe.fit(X_train,y_train)
13
14 y_pred = pipe.predict(X_test)
15
16 print('R2 score',r2_score(y_test,y_pred))
17 print('MAE',mean_absolute_error(y_test,y_pred))
```

R2 score 0.8073277448418643

MAE 0.2101782797642889

C:\Users\sanket parekh\miniconda3\lib\site-packages\sklearn\preprocessing_encoders.py:828: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.

```
warnings.warn(
```

Ridge Regression

In [89]:

```
1 step1 = ColumnTransformer(transformers=[  
2     ('col_tnf',OneHotEncoder(sparse=False,drop='first'),[0,1,7,10,11])  
3 ],remainder='passthrough')  
4  
5 step2 = Ridge(alpha=10)  
6  
7 pipe = Pipeline([  
8     ('step1',step1),  
9     ('step2',step2)  
10])  
11  
12 pipe.fit(X_train,y_train)  
13  
14 y_pred = pipe.predict(X_test)  
15  
16 print('R2 score',r2_score(y_test,y_pred))  
17 print('MAE',mean_absolute_error(y_test,y_pred))
```

R2 score 0.8127331031311809

MAE 0.20926802242582976

C:\Users\sanket parekh\miniconda3\lib\site-packages\sklearn\preprocessing_encoders.py:828: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.

```
    warnings.warn(
```

Lasso Regression

In [91]:

```
1 step1 = ColumnTransformer(transformers=[  
2     ('col_tnf',OneHotEncoder(sparse=False,drop='first'),[0,1,7,10,11])  
3 ],remainder='passthrough')  
4  
5 step2 = Lasso(alpha=0.001)  
6  
7 pipe = Pipeline([  
8     ('step1',step1),  
9     ('step2',step2)  
10])  
11  
12 pipe.fit(X_train,y_train)  
13  
14 y_pred = pipe.predict(X_test)  
15  
16 print('R2 score',r2_score(y_test,y_pred))  
17 print('MAE',mean_absolute_error(y_test,y_pred))  
18
```

R2 score 0.8071853945317103

MAE 0.21114361613472565

C:\Users\sanket parekh\miniconda3\lib\site-packages\sklearn\preprocessing_encoders.py:828: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.

```
    warnings.warn(
```

KNN

In [92]:

```
1 step1 = ColumnTransformer(transformers=[  
2     ('col_tnf',OneHotEncoder(sparse=False,drop='first'),[0,1,7,10,11])  
3 ],remainder='passthrough')  
4  
5 step2 = KNeighborsRegressor(n_neighbors=3)  
6  
7 pipe = Pipeline([  
8     ('step1',step1),  
9     ('step2',step2)  
10])  
11  
12 pipe.fit(X_train,y_train)  
13  
14 y_pred = pipe.predict(X_test)  
15  
16 print('R2 score',r2_score(y_test,y_pred))  
17 print('MAE',mean_absolute_error(y_test,y_pred))
```

C:\Users\sanket parekh\miniconda3\lib\site-packages\sklearn\preprocessing_encoders.py:828: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.
warnings.warn(

R2 score 0.803148868705085
MAE 0.19264883332948868

Decision Tree

In [93]:

```
1 step1 = ColumnTransformer(transformers=[  
2     ('col_tnf',OneHotEncoder(sparse=False,drop='first'),[0,1,7,10,11])  
3 ],remainder='passthrough')  
4  
5 step2 = DecisionTreeRegressor(max_depth=8)  
6  
7 pipe = Pipeline([  
8     ('step1',step1),  
9     ('step2',step2)  
10])  
11  
12 pipe.fit(X_train,y_train)  
13  
14 y_pred = pipe.predict(X_test)  
15  
16 print('R2 score',r2_score(y_test,y_pred))  
17 print('MAE',mean_absolute_error(y_test,y_pred))
```

R2 score 0.8319084067930679

MAE 0.18641541878794388

C:\Users\sanket parekh\miniconda3\lib\site-packages\sklearn\preprocessing_encoders.py:828: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.

```
    warnings.warn(
```

SVM

In [94]:

```
1 step1 = ColumnTransformer(transformers=[  
2     ('col_tnf',OneHotEncoder(sparse=False,drop='first'),[0,1,7,10,11])  
3 ],remainder='passthrough')  
4  
5 step2 = SVR(kernel='rbf',C=10000,epsilon=0.1)  
6  
7 pipe = Pipeline([  
8     ('step1',step1),  
9     ('step2',step2)  
10 ])  
11  
12 pipe.fit(X_train,y_train)  
13  
14 y_pred = pipe.predict(X_test)  
15  
16 print('R2 score',r2_score(y_test,y_pred))  
17 print('MAE',mean_absolute_error(y_test,y_pred))
```

C:\Users\sanket parekh\miniconda3\lib\site-packages\sklearn\preprocessing_encoders.py:828: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.
warnings.warn(

R2 score 0.8083180902289917
MAE 0.2023905942719158

Random Forest

In [95]:

```
1 step1 = ColumnTransformer(transformers=[  
2     ('col_tnf',OneHotEncoder(sparse=False,drop='first'),[0,1,7,10,11])  
3 ],remainder='passthrough')  
4  
5 step2 = RandomForestRegressor(n_estimators=100,  
6                             random_state=3,  
7                             max_samples=0.5,  
8                             max_features=0.75,  
9                             max_depth=15)  
10  
11 pipe = Pipeline([  
12     ('step1',step1),  
13     ('step2',step2)  
14 ])  
15  
16 pipe.fit(X_train,y_train)  
17  
18 y_pred = pipe.predict(X_test)  
19  
20 print('R2 score',r2_score(y_test,y_pred))  
21 print('MAE',mean_absolute_error(y_test,y_pred))  
22
```

C:\Users\sanket parekh\miniconda3\lib\site-packages\sklearn\preprocessing_encoders.py:828: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.
warnings.warn(

R2 score 0.8873402378382488
MAE 0.15860130110457718

Ada Boost

In [96]:

```
1 step1 = ColumnTransformer(transformers=[  
2     ('col_tnf',OneHotEncoder(sparse=False,drop='first'),[0,1,7,10,11])  
3 ],remainder='passthrough')  
4  
5 step2 = AdaBoostRegressor(n_estimators=15,learning_rate=1.0)  
6  
7 pipe = Pipeline([  
8     ('step1',step1),  
9     ('step2',step2)  
10])  
11  
12 pipe.fit(X_train,y_train)  
13  
14 y_pred = pipe.predict(X_test)  
15  
16 print('R2 score',r2_score(y_test,y_pred))  
17 print('MAE',mean_absolute_error(y_test,y_pred))
```

R2 score 0.8045879264218471

MAE 0.22466879624872899

C:\Users\sanket parekh\miniconda3\lib\site-packages\sklearn\preprocessing_encoders.py:828: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.

```
    warnings.warn(
```

Gradient Boost

In [97]:

```
1 step1 = ColumnTransformer(transformers=[  
2     ('col_tnf',OneHotEncoder(sparse=False,drop='first'),[0,1,7,10,11])  
3 ],remainder='passthrough')  
4  
5 step2 = GradientBoostingRegressor(n_estimators=500)  
6  
7 pipe = Pipeline([  
8     ('step1',step1),  
9     ('step2',step2)  
10 ])  
11  
12 pipe.fit(X_train,y_train)  
13  
14 y_pred = pipe.predict(X_test)  
15  
16 print('R2 score',r2_score(y_test,y_pred))  
17 print('MAE',mean_absolute_error(y_test,y_pred))
```

C:\Users\sanket parekh\miniconda3\lib\site-packages\sklearn\preprocessing_encoders.py:828: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.
warnings.warn(

R2 score 0.8810856073120187
MAE 0.15981480246214838

Xg Boost

```
In [98]: 1 step1 = ColumnTransformer(transformers=[  
2     ('col_tnf',OneHotEncoder(sparse=False,drop='first'),[0,1,7,10,11])  
3 ],remainder='passthrough')  
4  
5 step2 = XGBRegressor(n_estimators=45,max_depth=5,learning_rate=0.5)  
6  
7 pipe = Pipeline([  
8     ('step1',step1),  
9     ('step2',step2)  
10 ])  
11  
12 pipe.fit(X_train,y_train)  
13  
14 y_pred = pipe.predict(X_test)  
15  
16 print('R2 score',r2_score(y_test,y_pred))  
17 print('MAE',mean_absolute_error(y_test,y_pred))  
18
```

R2 score 0.8811773435850243

MAE 0.16496203512600974

C:\Users\sanket parekh\miniconda3\lib\site-packages\sklearn\preprocessing_encoders.py:828: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.

```
    warnings.warn(
```

In []:

1